

# Data all over the place!

How SQL and Apache Calcite bring  
sanity to streaming and spatial  
data

Julian Hyde | **Lyft** | 2018/06/27



# @julianhyde

---

SQL

Query planning

Query federation

OLAP

Streaming


Hadoop

ASF member

Original author of Apache Calcite

PMC Apache Arrow, Calcite, Drill, Eagle, Kylin

Architect at Looker



# Overview

---

Apache Calcite - What is it, exactly?

Streaming SQL / standardizing / how Calcite helps

Spatial / accelerated by materialized views / spatial + streaming

# 1. Calcite

---

# Apache Calcite



Apache top-level project


Query planning framework used in many projects and products

Also works standalone: embedded federated query engine with SQL / JDBC front end

Apache community development model

<http://calcite.apache.org>

<http://github.com/apache/calcite>



# Apache Calcite - Project goals

---

Make it easier to write a simple DBMS

Advance the state of the art for complex DBMS by pooling resources


Bring database approaches to new areas (e.g. streaming)

Allow create a DBMS by composing pieces (federation, etc.)

Customize by plugging into framework, evolving framework when necessary

Apache license & governance

# Architecture



**Core** – Operator expressions (relational algebra) and planner (based on Volcano/Cascades[2])


**External** – Data storage, algorithms and catalog

**Optional** – SQL parser, JDBC & ODBC drivers

**Extensible** – Planner rewrite rules, statistics, cost model, algebra, UDFs

# Planning queries


```
select p.productName, count(*) as c
from splunk.splunk as s
      join mysql.products as p
      on s.productId = p.productId
where s.action = 'purchase'
group by p.productName
order by c desc
```





# Optimized query

```
select p.productName, count(*) as c
from splunk.splunk as s
      join mysql.products as p
      on s.productId = p.productId
where s.action = 'purchase'
group by p.productName
order by c desc
```



# Calcite framework

## Relational algebra

RelNode (operator)

- TableScan
- Filter
- Project
- Union
- Aggregate
- ...

RelDataType (type)

RexNode (expression)

RelTrait (physical property)

- RelConvention (calling-convention)
- RelCollation (sortedness)
- RelDistribution (partitioning)

RelBuilder

## SQL parser

SqlNode

SqlParser

SqlValidator

## Metadata

Schema

Table

Function

- TableFunction
- TableMacro

Lattice

## JDBC driver

## Transformation rules

RelOptRule

- FilterMergeRule
- AggregateUnionTransposeRule
- 100+ more

Global transformations

- Unification (materialized view)
- Column trimming
- De-correlation

## Cost, statistics


RelOptCost

RelOptCostFactory


RelMetadataProvider

- RelMdColumnUniqueness
- RelMdDistinctRowCount
- RelMdSelectivity


# Conventions




1. Plans start as logical nodes.



2. Assign each Scan its table's native convention.



3. Fire rules to propagate conventions to other nodes.




4. The best plan may use an **engine** not tied to any native format.

To implement, generate a **program** that calls out to **query1** and **query2**.

# Adapters

---



Conventions provide a uniform representation for hybrid queries

Like ordering and distribution, conventions are a **physical property** of nodes

Adapter =  
    schema factory (lists tables)  
+ convention  
+ rules to convert nodes to a convention

Example adapters: file, web, MongoDB, Druid, Drill, Spark, Flink

## 2. Streaming


---

# Why SQL?

Data in motion, data at rest - it's all just data

Stream / database duality

SQL is the best language ever created for data (because it's declarative)



# Calcite and streaming SQL

---



Building a streaming SQL standard via consensus

- Please! No more “SQL-like” languages!

Toolkit for building a streaming SQL engine

- Key technologies are open source (many are Apache projects)
- Use Calcite’s framework to build a streaming SQL parser/planner for your engine

Developing example queries, data sets, TCK

# Simple queries

---

```
select *  
from Products  
where unitPrice < 20
```

- Traditional (non-streaming)
- `Products` is a table
- Retrieves records from  $-\infty$  to now

```
select stream *  
from Orders  
where units > 1000
```

- Streaming
- `Orders` is a stream
- Retrieves records from now to  $+\infty$
- Query never terminates



# Stream-table duality

---

```
select *  
from Orders  
where units > 1000
```

```
select stream *  
from Orders  
where units > 1000
```

- Yes, you can use a stream as a table
- And you can use a table as a stream
- Actually, `Orders` is both
- Use the `stream` keyword
- Where to actually find the data? That's up to the system

# Hybrid query combines a stream with its own history

---

- `Orders` is used as both as stream and as “stream history” virtual table
- “Average order size over last year” should be maintained by the system, i.e. a materialized view

“Orders” used as a stream

“Orders” used as a “stream history” virtual table

```
select stream *  
from Orders as o  
where units > (  
  select avg(units)  
  from Orders as h  
  where h.productId = o.productId  
  and h.rowtime > o.rowtime  
    - interval '1' year)
```

# Other operations

---

Other relational operations make sense on streams (usually only if there is an implicit time bound).

Examples:

- **order by** – E.g. Each hour emit the top 10 selling products
- **union** – E.g. Merge streams of orders and shipments
- **insert, update, delete** – E.g. Continuously insert into an external table
- **exists, in** sub-queries – E.g. Show me shipments of products for which there has been no order in the last hour
- **view** – Expanded when query is parsed; zero runtime cost
- **match\_recognize** – Complex event processing (CEP)

# “Standard streaming SQL”

---

Consistent semantics over tables (if database contains both streams and tables)

Consistent semantics over streams. The replay principle:

*A streaming query produces the same result as the corresponding non-streaming query would if given the same data in a table.*

Consistent look and feel (e.g. same data types, keywords)

Some hard problems remain... mainly concerning time...

# Controlling when data is emitted

---

Early emission is the defining characteristic of a streaming query.

The `emit` clause is a SQL extension inspired by Apache Beam's "trigger" notion. (Still experimental... and evolving.)

A relational (non-streaming) query is just a query with the most conservative possible emission strategy.

```
select stream productId,  
       count(*) as c  
from Orders  
group by productId,  
       floor(rowtime to hour)  
emit at watermark,  
     early interval '2' minute,  
     late limit 1
```

```
select *  
from Orders  
emit when complete
```

# Event time

---

Example streams and query so far have used a `rowtime` column

- Is `rowtime` a system column, or can streams choose their own sort key?
- Can a stream have more than one sort key?
- Is `rowtime`'s type always `TIMESTAMP`?
- Is a stream totally sorted on `rowtime`? Alternatives:
  - K-sorted: Every row is within 100 rows of sorted
  - T-sorted: Every row is within 5 minutes of sorted
  - Bounded latency: Every row is within 10 minutes of wallclock time
  - Watermark: An upper-bound on rows' delay, allows early emit

# Event time in queries

We've sorted out event time in streams... what about derived streams (queries / views)?

Given a streaming query:

- What are its sort keys?
- What is the maximum delay of rows?
- Does the system refuse to run queries with unbounded delay?

SQL syntax to promote a query column to a sort key

```
select count(*)  
from Orders
```

```
select count(*)  
from Orders  
group by floor(  
    rowtime to hour)
```

```
select shipTime  
    as rowtime  
from Orders  
order by shipTime
```



# Models of relations

---

A **table** is a multi-set of records; contents vary over time, with some kind transactional consistency

A **temporal table** is like a table, but you can also query its contents at previous times; usually mapped to a base table with (start, end) columns

So, what is a stream?

1. A **stream** is an append-only table with an event-time column
2. Or, if you speak monads[3], a **stream** is a function:  $Stream\ x \rightarrow Time \rightarrow Bag\ x$
3. Or, a **stream** is the time-derivative of a table

Are a table, its temporal table, and its stream one catalog object or three?

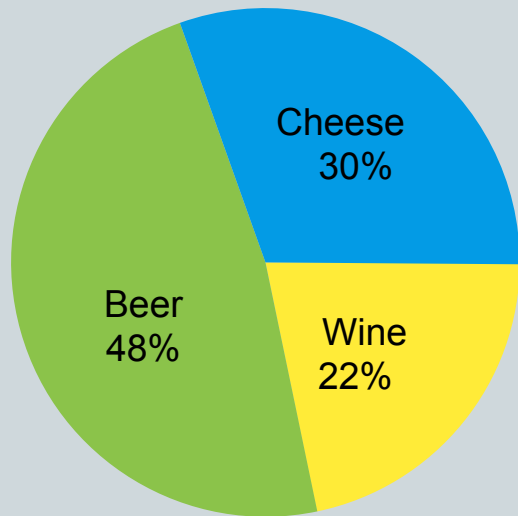


# The “pie chart” problem

- **Task** – Write dashboard of orders over last hour
- **Problem** – The *Orders* stream only contains the current few records
- **Solution** – Materialize short-term history
- **Question** – Does the developer maintain that mapping, or does the DBMS do it automatically?

```
select productId, count(*)  
from Orders  
where rowtime > current_timestamp - interval '1' hour  
group by productId
```

*Orders over the last hour*



# Transactions & isolation

---

In a regular database, INSERT, UPDATE and DELETE are transactional operations; SELECT is not

In a streaming database, SELECT might be transactional too

- Consider a pub-sub system: the system acknowledges when you send a message, you must acknowledge when you receive a message

In a regular database query, snapshot isolation freezes the database in time; in a streaming query, we don't want that

# Joins

---

Stream to stream join

```
select stream *  
from Orders as o  
join Shipments as s  
  on s.orderId = o.orderId  
  and s.rowtime between o.rowtime  
    and o.rowtime + interval '1' hour
```

Stream to table join

```
select stream *  
from Orders as o  
join Products as p  
  on p.productId = o.productId
```

Stream to temporal table join

```
select stream *  
from Orders as o  
join TemporalProducts as of o.rowtime  
  as p  
  on p.productId = o.productId
```


“as of” is based on SQL:2011 temporal syntax

# 3. Spatial

---

# Spatial query

Find all restaurants within 1.5 distance units of my location (6, 7)




restaurant	x	y
Zachary's pizza	3	1
King Yen	7	7
Filippo's	7	4
Station burger	5	6

# Spatial query

Find all restaurants within 1.5 distance units of my location (6, 7)

Using OpenGIS SQL extensions:

```
SELECT *  
FROM Restaurants AS r  
WHERE ST_Distance(  
    ST_MakePoint(r.x, r.y),  
    ST_MakePoint(6, 7)) < 1.5
```



restaurant	x	y
Zachary's pizza	3	1
King Yen	7	7
Filippo's	7	4
Station burger	5	6

# Efficient spatial query?

Spatial queries are hard to execute efficiently.


Conventional DBs use two tricks that don't work:

- Hash (no ranges / no locality)
- Sort (only one-dimensional ranges)

Two techniques that do work[4]:


- Data-organized structures (e.g. r-tree) – self-balancing
- Space-organized structures (e.g. tiles) – allow algebraic rewrite

Algebraic rewrites work on general-purpose data structures, and they compose.



A scan over a two-dimensional index only has locality in one dimension

# Hilbert space-filling curve



- A space-filling curve invented by mathematician David Hilbert
- Every  $(x, y)$  point has a unique position on the curve
- Points near to each other typically have Hilbert indexes close together




# Using Hilbert index

Add restriction based on **h**, a restaurant's distance along the Hilbert curve

Must keep original restriction due to false positives

```
SELECT *  
FROM Restaurants AS r  
WHERE (r.h BETWEEN 35 AND 42  
      OR r.h BETWEEN 46 AND 46)  
AND ST_Distance(  
  ST_MakePoint(r.x, r.y),  
  ST_MakePoint(6, 7)) < 1.5
```



restaurant	x	y	h
Zachary's pizza	3	1	5
King Yen	7	7	41
Filippo's	7	4	52
Station burger	5	6	36

# Telling the optimizer

1. Declare **h** as a generated column
2. Sort table by **h**

Planner can now convert spatial range queries into a range scan

Does not require specialized spatial index such as r-tree

Very efficient on a sorted table such as HBase

```
CREATE TABLE Restaurants (  
  restaurant VARCHAR(20),  
  x DOUBLE,  
  y DOUBLE,  
  h DOUBLE GENERATED ALWAYS AS  
    ST_Hilbert(x, y) STORED)  
SORT KEY (h);
```


restaurant	x	y	h
Zachary's pizza	3	1	5
Station burger	5	6	36
King Yen	7	7	41
Filippo's	7	4	52

# More spatial queries

---

What state **am I in?** (point-to-polygon)


Which states does **Yellowstone NP** intersect?  
(polygon-to-polygon)



# More spatial queries

What state **am I in?** (point-to-polygon)

Which states does **Yellowstone NP** intersect?  
(polygon-to-polygon)




```
SELECT *  
FROM States AS s  
WHERE ST_Intersects(s.geometry,  
    ST_MakePoint(6, 7))
```

```
SELECT *  
FROM States AS s  
WHERE ST_Intersects(s.geometry,  
    ST_GeomFromText('LINESTRING(...)' ))
```

# Point-to-polygon query

What state am I in? (point-to-polygon)

1. Divide the plane into tiles, and pre-compute the state-tile intersections
2. Use this 'tile index' to narrow list of states




```
SELECT s.*
FROM States AS s
WHERE s.stateId IN (SELECT stateId
    FROM StateTiles AS t
    WHERE t.tileId = 8)
AND ST_Intersects(s.geometry, ST_MakePoint(6, 7))
```

# Building the tile index

Use the `ST_MakeGrid` function to decompose each state into a series of tiles

Store the results in a table, sorted by tile id

A materialized view is a table that remembers how it was computed, so the planner can rewrite queries to use it



```
CREATE MATERIALIZED VIEW StateTiles AS
SELECT s.stateId, t.geometry, t.tileId, t.id_col, t.id_row
FROM States AS s,
     LATERAL TABLE(ST_MakeGrid(s.stateId, 4, 4)) AS t
```

# Streaming + spatial

---

Every minute, emit the number of journeys that have intersected each city. (Some journeys intersect multiple cities.)

(Efficient implementation is left as an exercise to the reader. Probably involves splitting journeys into tiles, partitioning by tile hash-code, intersecting with cities in those tiles, then rolling up cities.)

```
SELECT STREAM c.name, COUNT(*)  
FROM Journeys AS j  
CROSS JOIN Cities AS c  
  ON ST_Intersects(c.geometry, j.geometry)  
GROUP BY c.name, FLOOR(j.rowtime TO HOUR)
```

# 4. Conclusion

---



# Why SQL?


---

Common language across paradigms (OLTP, analytics, streaming, spatial, graph, search...) and workloads (batch, transactional, continuous)

SQL is standard (and growing)

Relational algebra allows a planner to transform queries

Materialized views (indexes, caches) allow you to optimize your data



# Thank you! Questions?

## References

1. Edmon Begoli, Jesús Camacho Rodríguez, Julian Hyde, Michael Mior, Daniel Lemire. 2018. Apache Calcite: A Foundational Framework for Optimized Query Processing Over Heterogeneous Data Sources. [SIGMOD 2018](#).
2. Goetz Graefe and William J. McKenna. 1993. The Volcano Optimizer Generator: Extensibility and Efficient Search. [ICDE](#).
3. Sebastian Herbst, Andreas M. Wahl, Johannes Tenschert, Klaus Meyer-Wegener. 2018. Sequences, yet Functions: The Dual Nature of Data-Stream Processing. [arXiv preprint](#).
4. Edwin H. Jacox, Hanan Samet. 2007. Spatial join techniques. [TODS](#).

<https://calcite.apache.org>  
@ApacheCalcite  
@julianhyde





# Extra slides

---

# Talk proposal

---

## Bio

Julian Hyde is an expert in query optimization and in-memory analytics. He founded Apache Calcite, a framework for query optimization and data virtualization. He also founded Mondrian, the popular open source OLAP engine, and co-founded SQLstream, an early streaming SQL platform. He is an architect at Looker.

## Title

Data all over the place! How Apache Calcite brings SQL and sanity to streaming and spatial data

## Abstract

The revolution has happened. We are living the age of the deconstructed database. The modern enterprises are powered by data, and that data lives in many formats and locations, in-flight and at rest, but somewhat surprisingly, the lingua franca for remains SQL. In this talk, Julian describes Apache Calcite, a toolkit for relational algebra that powers many systems including Apache Beam, Flink and Hive. He discusses some areas of development in Calcite: streaming SQL, materialized views, enabling spatial query on vanilla databases, and what a mash-up of all three might look like.

# Other applications of data profiling

---

Query optimization:

- Planners are poor at estimating selectivity of conditions after N-way join (especially on real data)
- New join-order benchmark: “Movies made by French directors tend to have French actors”
- Predict number of reducers in MapReduce & Spark

“Grokking” a data set

Identifying problems in normalization, partitioning, quality

Applications in machine learning?

# Further improvements

---

- Build sketches in parallel
- Run algorithm in a distributed framework (Spark or MapReduce)
- Compute histograms
  - For example, Median age for male/female customers
- Seek out functional dependencies
  - Once you know FDs, a lot of cardinalities are no longer “surprising”
  - FDs occur in denormalized tables, e.g. star schemas
- Smarter criteria for stopping algorithm
- Skew/heavy hitters. Are some values much more frequent than others?
- Conditional cardinalities and functional dependencies
  - Does one partition of the data behave differently from others? (e.g. year=2005, state=LA)

# Relational algebra (plus streaming)

---

## Core operators:

- Scan
- Filter
- Project
- Join
- Sort
- Aggregate
- Union
- Values

## Streaming operators:

- Delta (converts relation to stream)
- Chi (converts stream to relation)


In SQL, the STREAM keyword signifies Delta



# Streaming algebra

---

- Filter
- Route
- Partition
- Round-robin
- Queue
- Aggregate
- Merge
- Store
- Replay
- Sort
- Lookup



# Optimizing streaming queries

---

The usual relational transformations still apply: push filters and projects towards sources, eliminate empty inputs, etc.


The transformations for delta are mostly simple:

- $\text{Delta}(\text{Filter}(r, \text{predicate})) \rightarrow \text{Filter}(\text{Delta}(r), \text{predicate})$
- $\text{Delta}(\text{Project}(r, e_0, \dots)) \rightarrow \text{Project}(\text{Delta}(r), e_0, \dots)$
- $\text{Delta}(\text{Union}(r_0, r_1), \text{ALL}) \rightarrow \text{Union}(\text{Delta}(r_0), \text{Delta}(r_1))$

But not always:

- $\text{Delta}(\text{Join}(r_0, r_1, \text{predicate})) \rightarrow \text{Union}(\text{Join}(r_0, \text{Delta}(r_1)), \text{Join}(\text{Delta}(r_0), r_1))$
- $\text{Delta}(\text{Scan}(a\text{Table})) \rightarrow \text{Empty}$

# Sliding windows in SQL



```
select stream
  sum(units) over w (partition by productId) as units1hp,
  sum(units) over w as units1h,
  rowtime, productId, units
from Orders
window w as (order by rowtime range interval '1' hour preceding)
```

rowtime	productId	units		units1hp	units1h	rowtime	productId	units
09:12	100	5	→	5	5	09:12	100	5
09:25	130	10	→	10	15	09:25	130	10
09:59	100	3	→	8	18	09:59	100	3
10:17	100	10	→	23	13	10:17	100	10

# Join stream to a *changing* table

---

Execution is more difficult if the **Products** table is being changed while the query executes.

To do things properly (e.g. to get the same results when we re-play the data), we'd need temporal database semantics.

(Sometimes doing things properly is too expensive.)

```
select stream *  
from Orders as o  
join Products as p  
  on o.productId = p.productId  
  and o.rowtime  
    between p.startEffectiveDate  
    and p.endEffectiveDate
```

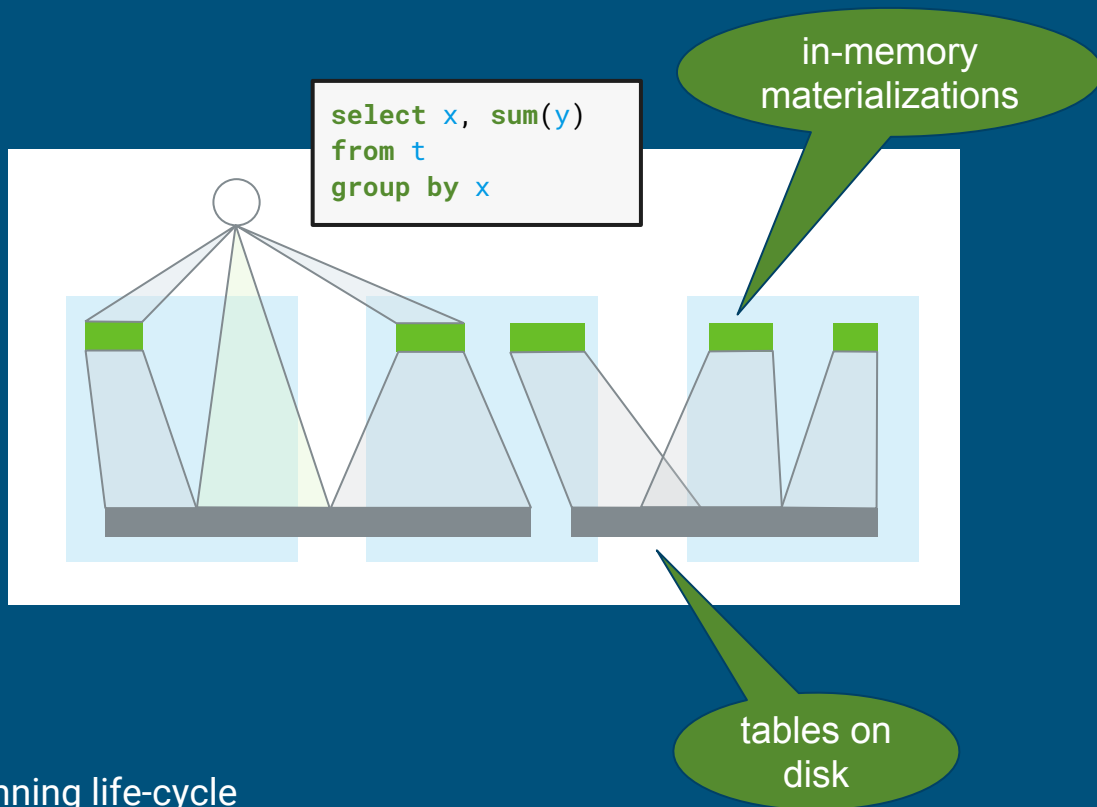
# Calcite design

## Design goals:

- Not-just-SQL front end
- Federation
- Extensibility
- Caching / hybrid storage
- Materialized views

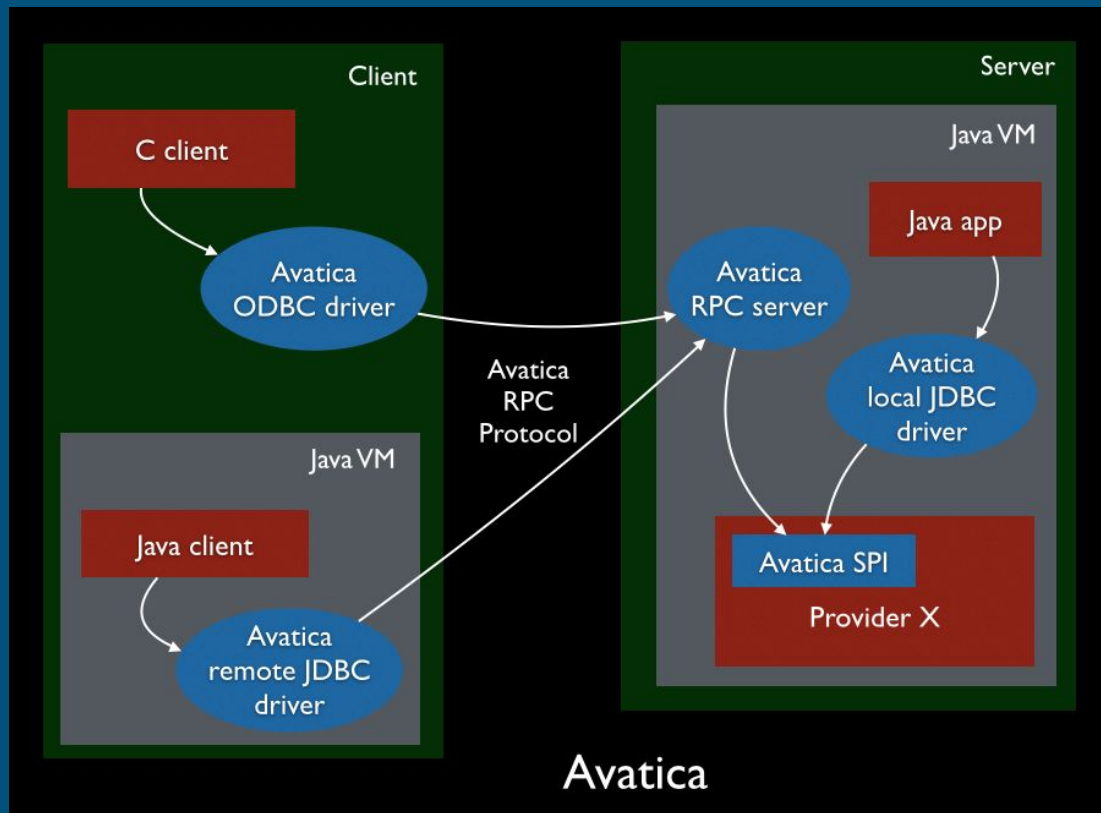
## Design points:

- Relational algebra foundation
- Composable transformation rules
- Support different physical engines
- Pluggable cost model, statistics, planning life-cycle



# Avatica

- Database connectivity stack
- Self-contained sub-project of Calcite
- Fast, open, stable
- Protobuf or JSON over HTTP
- Powers Phoenix Query Server






# Aggregation and windows on streams

**GROUP BY** aggregates multiple rows into sub-totals


- In regular GROUP BY each row contributes to exactly one sub-total
- In multi-GROUP BY (e.g. HOP, GROUPING SETS) a row can contribute to more than one sub-total

**Window functions** (OVER) leave the number of rows unchanged, but compute extra expressions for each row (based on neighboring rows)



# Tumbling, hopping & session windows in SQL


Tumbling window



```
select stream ... from Orders  
group by floor(rowtime to hour)
```


```
select stream ... from Orders  
group by tumble(rowtime, interval '1' hour)
```

Hopping window



```
select stream ... from Orders  
group by hop(rowtime, interval '1' hour,  
             interval '2' hour)
```

Session window



```
select stream ... from Orders  
group by session(rowtime, interval '1' hour)
```