

Smarter Together

**Bringing Relational Algebra,
Powered by Apache Calcite, into
Looker's Query Engine**

[Julian Hyde | JOIN 2019 | 5-7 November 2019 | San Francisco]



An algebra problem

Ice cream costs \$1.50

Topping costs \$0.50

How much do 3 ice creams with toppings cost?



An algebra problem

How did you solve it?

Ice cream costs \$1.50

Topping costs \$0.50

How much do 3 ice creams with toppings cost?



Method 1 $(3 * \$1.50) + (3 * \$0.50) = \$4.50 + \$1.50 = \$6.00$

Method 2 $3 * (\$1.50 + \$0.50) = 3 * \$2.00 = \6.00

An algebra problem

It's best to apply toppings first

Ice cream costs \$1.50

Topping costs \$0.50

How much do 3 ice creams with toppings cost?



Method 1 $(3 * \$1.50) + (3 * \$0.50) = \$4.50 + \$1.50 = \$6.00$

Method 2 $3 * (\$1.50 + \$0.50) = 3 * \$2.00 = \6.00

Algebraic identity: $(a * b) + (a * c) = a * (b + c)$

Algebra

Algebra is a mathematical language for combining values using operators.

Algebraic identities tell us how the operators and values interact.

We can prove things by using the identities as **rewrite rules**.

Values numbers

Operators $+$ $-$ $*$ $/$

Algebraic identities

- $a * (b + c) = (a * b) + (a * c)$
- $a + 0 = a$
- $a * 0 = 0$
- $a + b = b + a$

Relational algebra

Relational algebra is an algebra where the values are relations (aka tables, multisets of records).

Values relations

Operators join (\bowtie) union (\cup)
project (Π) filter (σ)

Algebraic identities

- $A \bowtie (B \cup C) = A \bowtie B \cup A \bowtie C$
- $A \cup \emptyset = A$
- $A \bowtie \emptyset = A$
- $A \cup B = B \cup A$
- $\sigma_P(\sigma_Q(A)) = \sigma_{P \wedge Q}(A)$

Algebraic identities in SQL

$$A \bowtie (B \cup C) = A \bowtie B \cup A \bowtie C$$

```
SELECT * FROM products
JOIN (SELECT FROM orders_2018
      UNION ALL
      SELECT * FROM orders_2019)
```

=

```
SELECT * FROM products JOIN orders_2018
UNION ALL
SELECT * FROM products JOIN orders_2019
```

$$\sigma_P(\sigma_Q(A)) = \sigma_{P \wedge Q}(A)$$

```
SELECT *
FROM (SELECT * FROM products
      WHERE color = 'red')
WHERE type = 'bicycle'
```

=

```
SELECT * FROM products
WHERE color = 'red'
AND type = 'bicycle'
```

Agenda

Relational algebra and query optimization

How Looker executes a query (and how Calcite can help)

Aggregate awareness & materialized views

Generate LookML model from SQL

Agenda

Relational algebra and query optimization

How Looker executes a query (and how Calcite can help)

Aggregate awareness & materialized views

Generate LookML model from SQL

Apache Calcite



Apache top-level project

Query planning framework used in many projects and products

Also works standalone: embedded federated query engine with SQL / JDBC front end

Apache community development model

<https://calcite.apache.org>

<https://github.com/apache/calcite>

Used by



Connects to



Apache Calcite – goals

Make it easier to write a simple DBMS

Advance the state of the art for complex DBMS by pooling resources

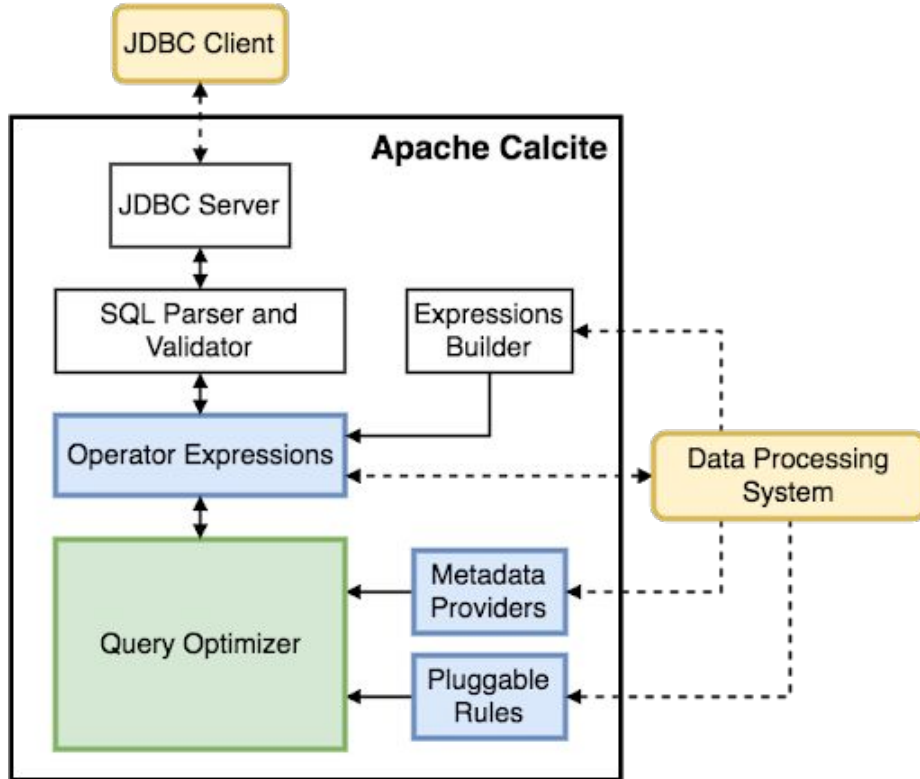
Bring database approaches to new areas (e.g. streaming)

Allow create a DBMS by composing pieces (federation, etc.)

Customize by plugging into framework, evolving framework when necessary

Apache license & governance

Apache Calcite - architecture



Core – Operator expressions (relational algebra) and planner (based on Volcano/Cascades)

External – Data storage, algorithms and catalog

Optional – SQL parser, JDBC & ODBC drivers

Extensible – Planner rewrite rules, statistics, cost model, algebra, UDFs

Relational algebra

Translation to SQL

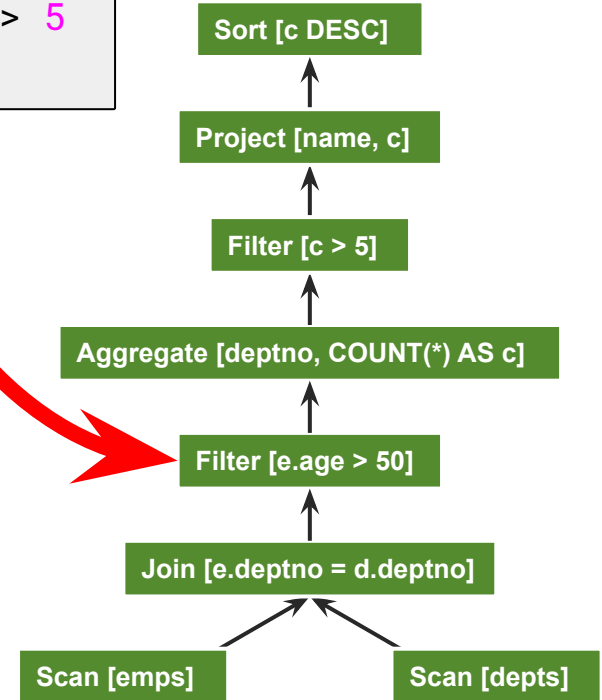
Based on set theory, plus
operators: Project, Filter,
Aggregate, Union, Join, Sort

Requires: declarative language
(SQL), query planner

Original goal: data independence

Enables: query optimization, new
algorithms and data structures

```
SELECT d.name, COUNT(*) AS c
FROM emps AS e
JOIN depts AS d USING (deptno)
WHERE e.age > 50
GROUP BY d.deptno
HAVING COUNT(*) > 5
ORDER BY c DESC
```



Relational algebra

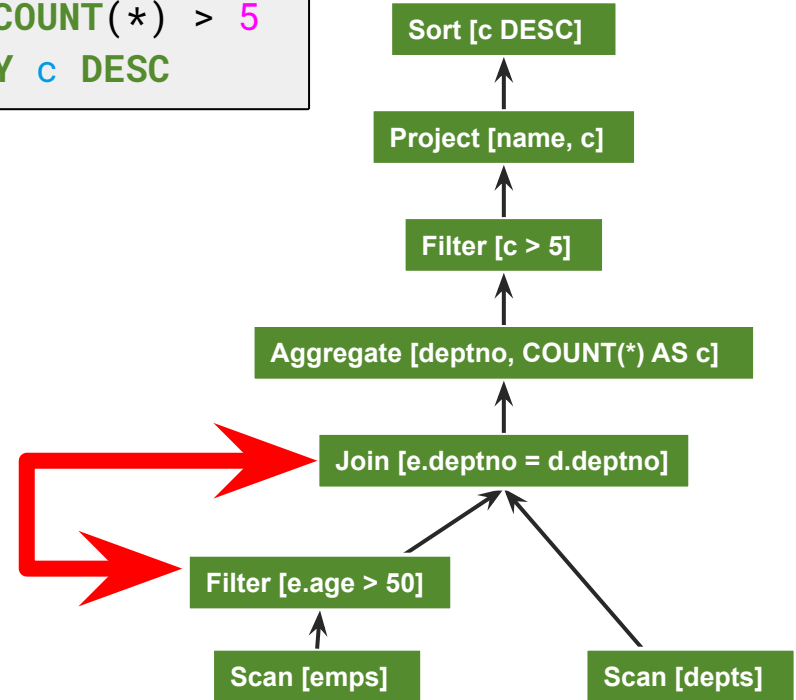
Algebraic rewrite

Optimize by applying rewrite rules that preserve semantics

Hopefully the result is less expensive; but it's OK if it's not (planner keeps "before" and "after")

Planner uses dynamic programming, seeking the lowest total cost

```
SELECT d.name, COUNT(*) AS c
FROM emps AS e
JOIN depts AS d USING (deptno)
WHERE e.age > 50
GROUP BY d.deptno
HAVING COUNT(*) > 5
ORDER BY c DESC
```



Calcite framework

Relational algebra

RelNode (operator)

- TableScan
- Filter
- Project
- Union
- Aggregate
- ...

RelDataType (type)

RexNode (expression)

RelTrait (physical property)

- RelConvention (calling-convention)
- RelCollation (sortedness)
- RelDistribution (partitioning)

RelBuilder

SQL parser

SqlNode

SqlParser

SqlValidator

Metadata

Schema

Table

Function

- TableFunction
- TableMacro

Lattice

JDBC driver

Remote driver

Local driver

Transformation rules

RelOptRule

- FilterMergeRule
- AggregateUnionTransposeRule
- 100+ more

Global transformations

- Unification (materialized view)
- Column trimming
- De-correlation

Cost, statistics

RelOptCost

RelOptCostFactory

RelMetadataProvider

- RelMdColumnUniqueness
- RelMdDistinctRowCount
- RelMdSelectivity

Agenda

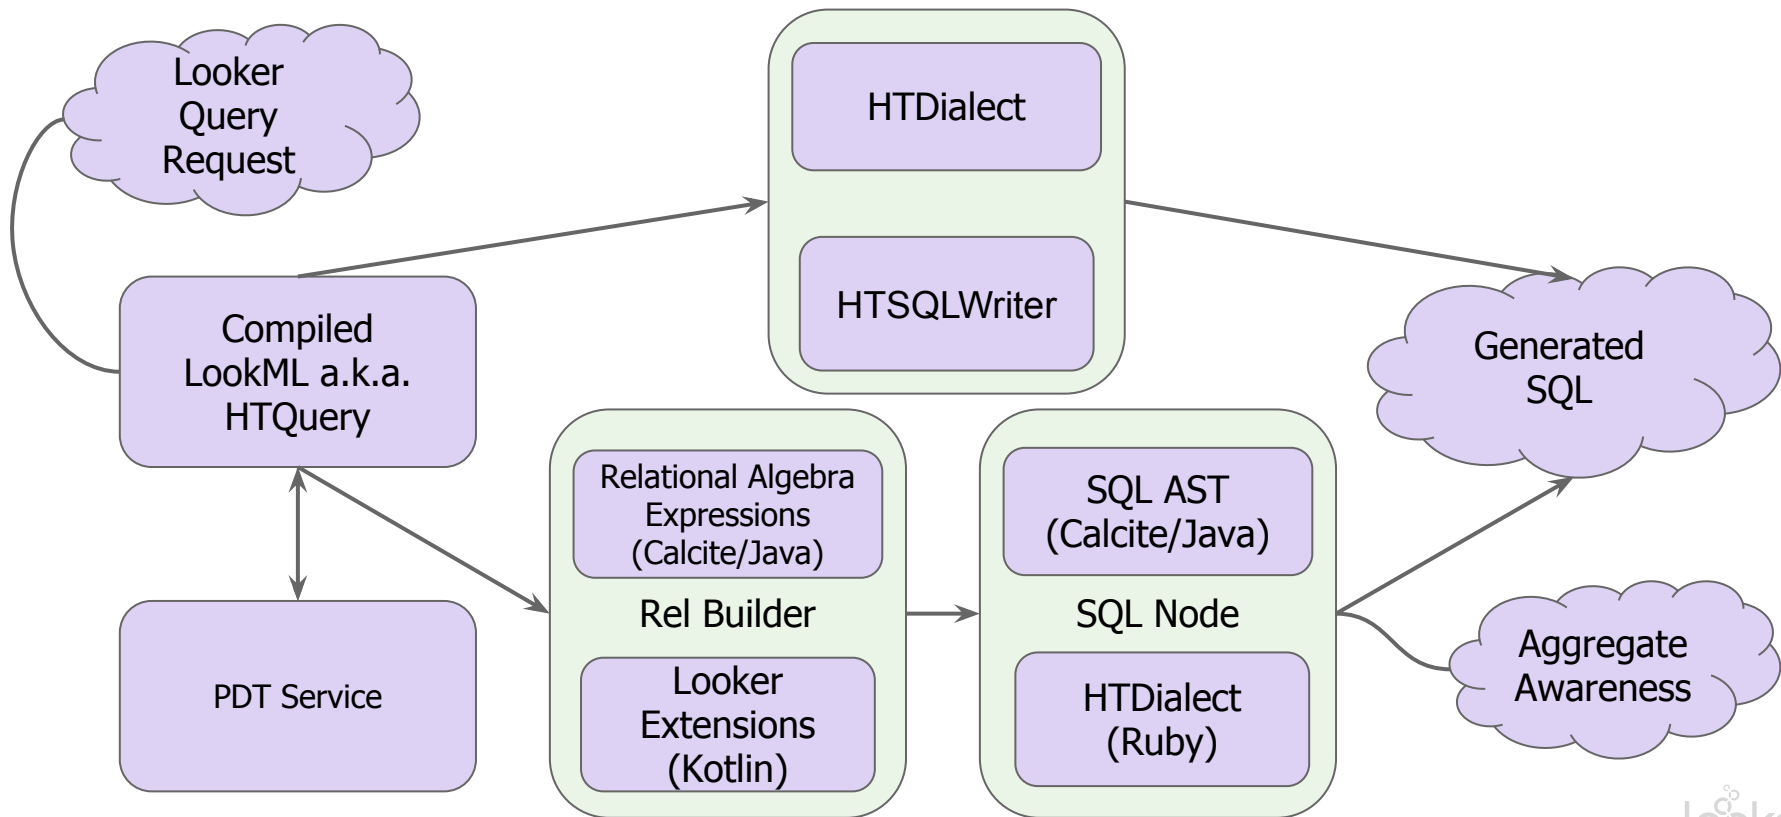
Relational algebra and query optimization

How Looker executes a query (and how Calcite can help)

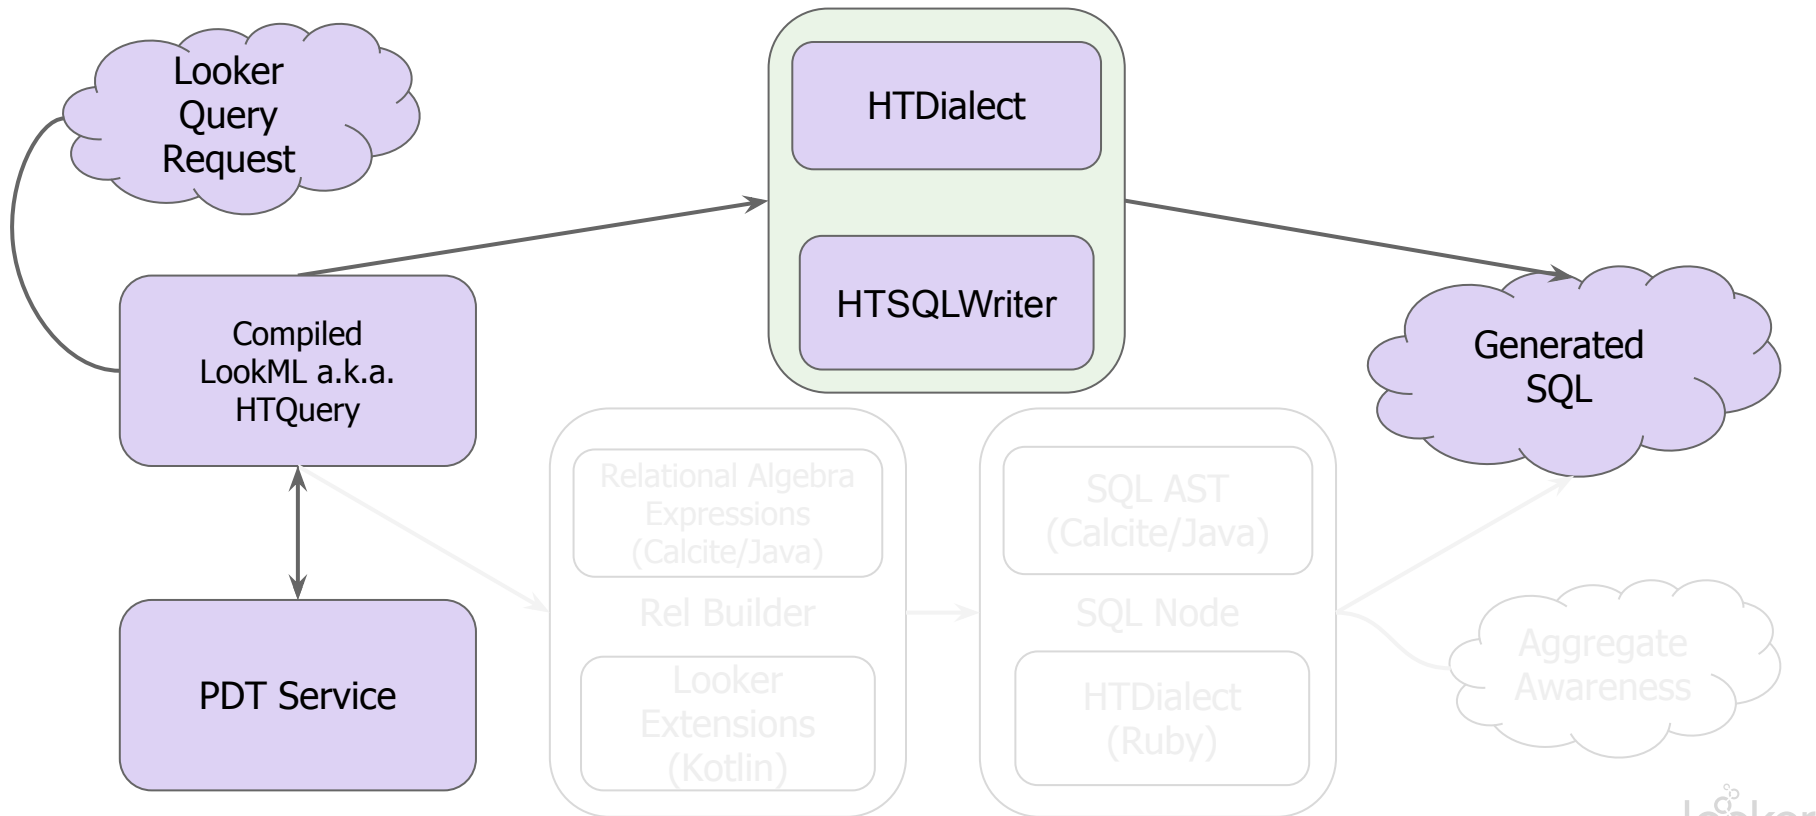
Aggregate awareness & materialized views

Generate LookML model from SQL

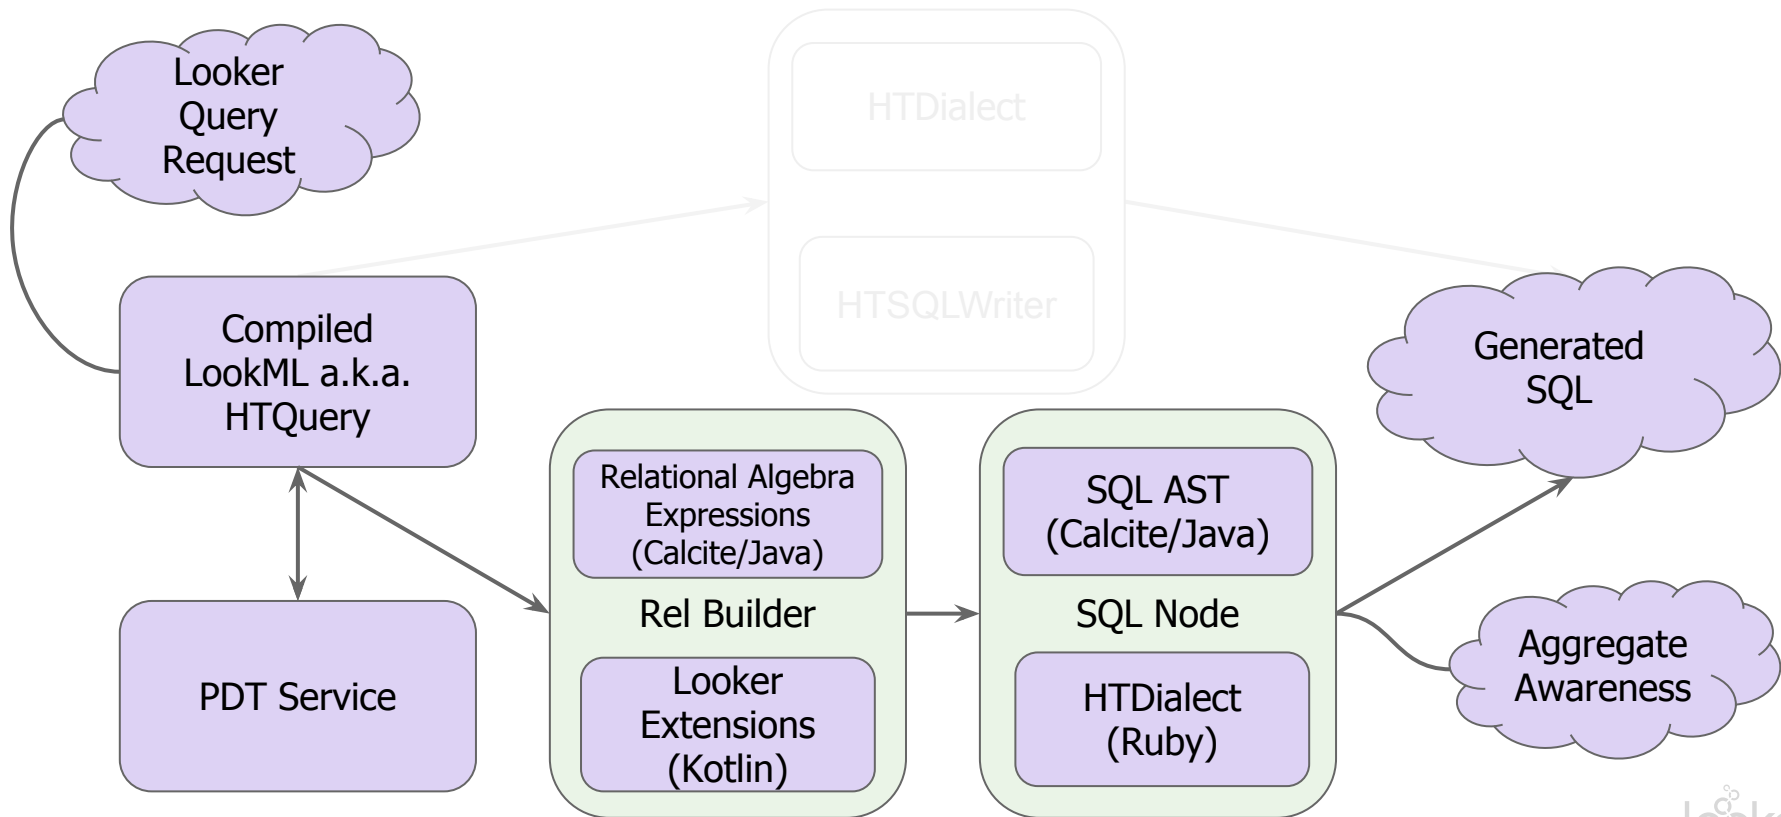
Looker SQL Generation Overview



Looker SQL Generation (old)



Looker SQL Generation (new)



Agenda

Relational algebra and query optimization

How Looker executes a query (and how Calcite can help)

Aggregate awareness & materialized views

Generate LookML model from SQL

Aggregates to the rescue

Business question:

**Which state had the
most revenue in 2018?**

customers (12 million rows)

customer_id	state
12345	CA
23456	TX



orders (100 million rows)

date	product_id	customer_id	revenue
2018-01-05	100	12345	\$17
2018-01-07	105	23456	\$120

Aggregates to the rescue

Business question:

**Which state had the
most revenue in 2018?**

Raw data – 120 million rows

Summary data – 2,000 rows

customers (12 million rows)

customer_id	state
12345	CA
23456	TX



orders (100 million rows)

date	product_id	customer_id	revenue
2018-01-05	100	12345	\$17
2018-01-07	105	23456	\$120

orders_by_state_and_quarter (2,000 rows)

quarter	state	sum_revenue	sales_count
2018-Q1	CA	\$15,567	46
2018-Q1	TX	\$23,006	88

Aggregates in LookML

LookML - orders

```
view: orders {  
  dimension_group: created_at {  
    type: time  
    sql: ${TABLE}.created_at ;;  
  }  
  dimension: product_id {}  
  dimension: customer_id {}  
  measure: revenue {}  
}  
  
view: customers {  
  dimension: customer_id {}  
  dimension: state {}  
}  
  
explore: orders {  
  join: customers {}  
}
```

LookML - orders_by_state_quarter

```
view: orders_by_state_quarter {  
  derived_table: {  
    explore_source: orders {  
      column: created_at_quarter {  
        field: orders.created_at_quarter  
      }  
      column: state {  
        field: customers.state  
      }  
      column: sum_revenue {  
        sql: sum(${orders.revenue});;  
      }  
    }  
  }  
  explore: orders_by_state_quarter {}  
}
```

NOTE: This feature is not production, and syntax may change

Views, materialized views, aggregate awareness

A **view** is a virtual table that expands to a relational expression

- No data stored on disk (except its definition)
- Purpose: abstraction

A **materialized view** is a table whose contents are defined by a relational expression

- Purpose: improve query performance

A **aggregate table** or **summary table** is a materialized view that uses GROUP BY (and perhaps JOIN) and has many fewer rows than its base table

A system with **aggregate awareness** automatically builds aggregate tables and uses them to make queries faster

View

```
SELECT deptno, MIN(salary)
FROM managers
WHERE age > 50
GROUP BY deptno
```

Query that uses the
managers view

```
CREATE VIEW managers AS
SELECT *
FROM emps AS e
WHERE EXISTS (
    SELECT *
    FROM emps AS underling
    WHERE underling.manager = e.id)
```

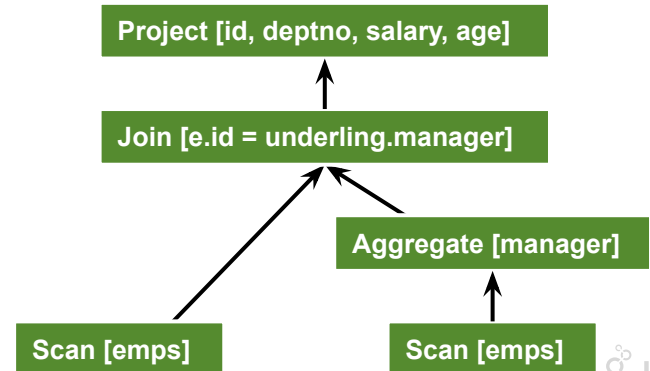

Declaration of the
managers view

View

With relational algebra

```
SELECT deptno, MIN(salary)
FROM managers
WHERE age > 50
GROUP BY deptno
```

```
CREATE VIEW managers AS
SELECT *
FROM emps AS e
WHERE EXISTS (
  SELECT *
  FROM emps AS underling
  WHERE underling.manager = e.id)
```



View

Showing the view and the relational expression that will replace it

```
SELECT deptno, MIN(salary)
FROM managers
WHERE age > 50
GROUP BY deptno
```

```
CREATE VIEW managers AS
SELECT *
FROM emps AS e
WHERE EXISTS (
  SELECT *
  FROM emps AS underling
  WHERE underling.manager = e.id)
```

Aggregate [deptno, MIN(salary)]

Filter [age > 50]

Scan [managers]

Project [id, deptno, salary, age]

Join [e.id = underling.manager]

Aggregate [manager]

Scan [emps]

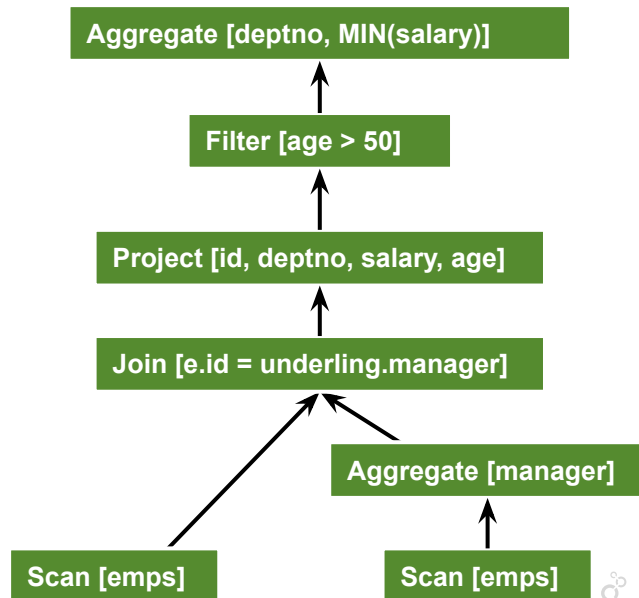
Scan [emps]

View

After view has been expanded

```
SELECT deptno, MIN(salary)
FROM managers
WHERE age > 50
GROUP BY deptno
```

```
CREATE VIEW managers AS
SELECT *
FROM emps AS e
WHERE EXISTS (
  SELECT *
  FROM emps AS underling
  WHERE underling.manager = e.id)
```

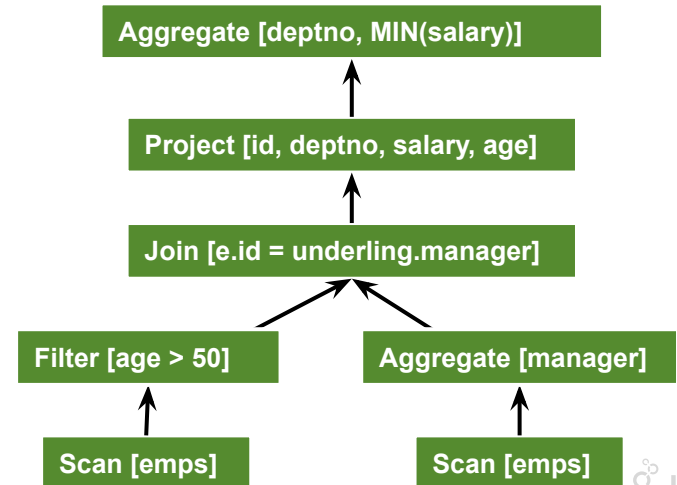


View

After pushing down “age > 50” filter

```
SELECT deptno, MIN(salary)
FROM managers
WHERE age > 50
GROUP BY deptno
```

```
CREATE VIEW managers AS
SELECT *
FROM emps AS e
WHERE EXISTS (
  SELECT *
  FROM emps AS underling
  WHERE underling.manager = e.id)
```



Materialized view

```
CREATE MATERIALIZED VIEW  
  emps_by_deptno_gender AS  
SELECT deptno, gender,  
       COUNT(*) AS c, SUM(sal) AS s  
FROM emps  
GROUP BY deptno, gender
```

Declaration of the
`emps_by_deptno_gender`
materialized view

```
SELECT COUNT(*) AS c  
FROM emps  
WHERE deptno = 10  
AND gender = 'M'
```

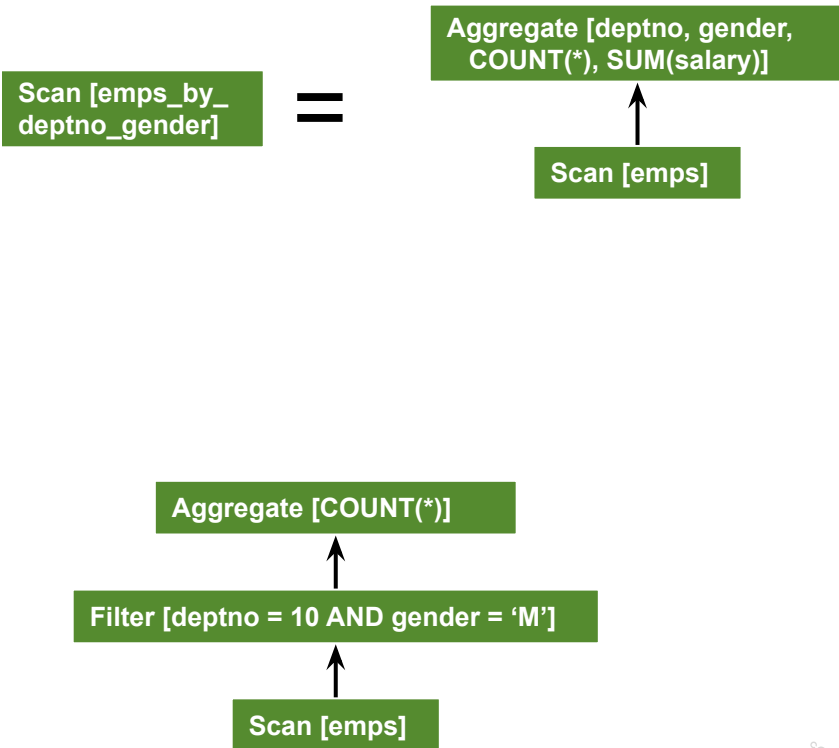
Query that uses the `emps` table
but could potentially use the
`emps_by_deptno_gender`
materialized view

Materialized view

With relational algebra

```
CREATE MATERIALIZED VIEW  
  emps_by_deptno_gender AS  
SELECT deptno, gender,  
       COUNT(*) AS c, SUM(sal) AS s  
FROM emps  
GROUP BY deptno, gender
```

```
SELECT COUNT(*) AS c  
FROM emps  
WHERE deptno = 10  
AND gender = 'M'
```

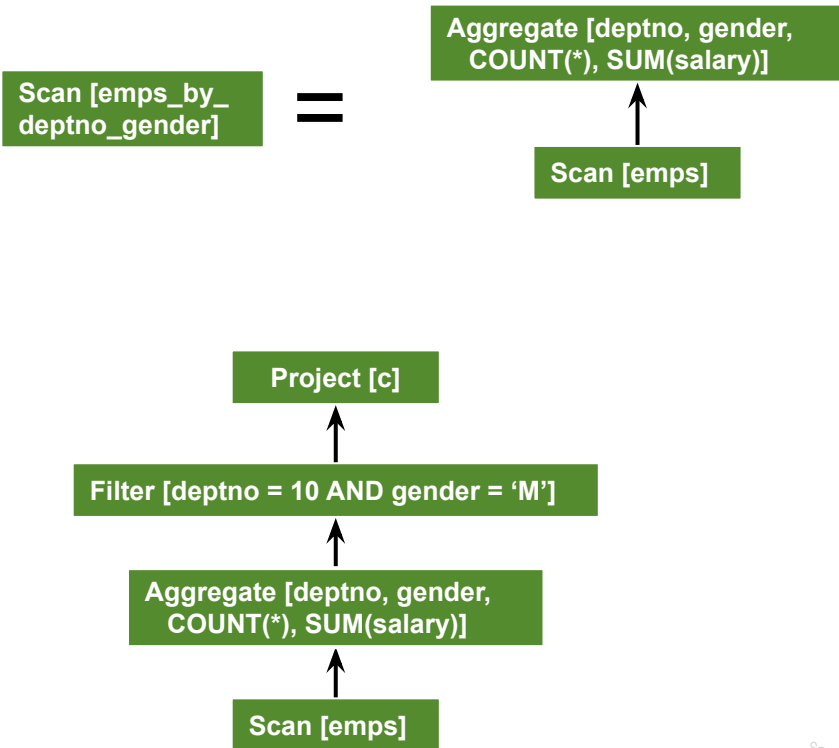


Materialized view

Rewrite query to match

```
CREATE MATERIALIZED VIEW  
  emps_by_deptno_gender AS  
SELECT deptno, gender,  
       COUNT(*) AS c, SUM(sal) AS s  
FROM emps  
GROUP BY deptno, gender
```

```
SELECT COUNT(*) AS c  
FROM emps  
WHERE deptno = 10  
AND gender = 'M'
```



Materialized view

Fragment of query is now identical to materialized view

```
CREATE MATERIALIZED VIEW  
  emps_by_deptno_gender AS  
SELECT deptno, gender,  
       COUNT(*) AS c, SUM(sal) AS s  
FROM emps  
GROUP BY deptno, gender
```

```
SELECT COUNT(*) AS c  
FROM emps  
WHERE deptno = 10  
AND gender = 'M'
```

Scan [emps_by_
deptno_gender]

=

Aggregate [deptno, gender,
COUNT(*), SUM(salary)]



Scan [emps]

Project [c]



Filter [deptno = 10 AND gender = 'M']



Aggregate [deptno, gender,
COUNT(*), SUM(salary)]



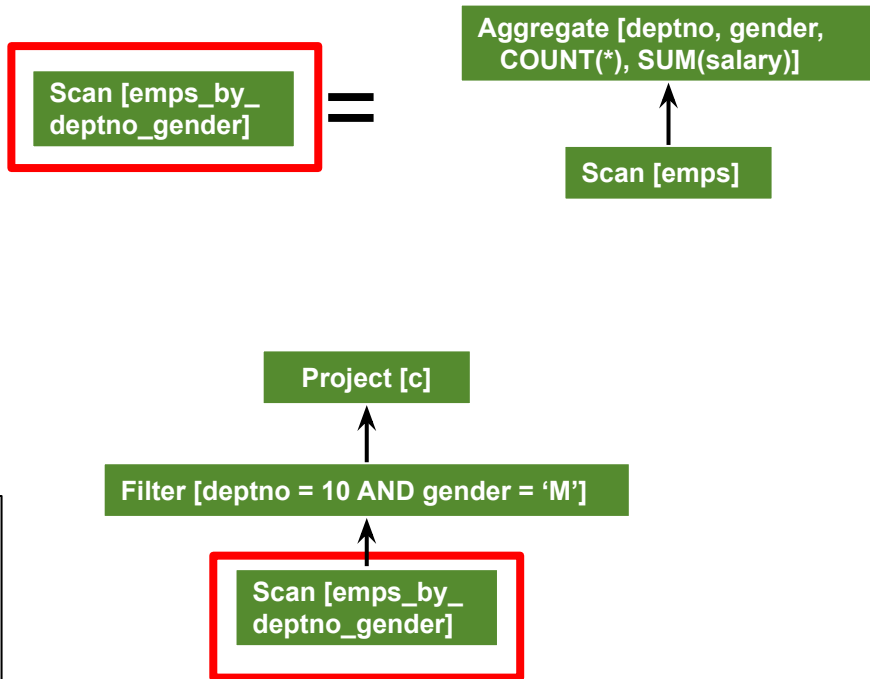
Scan [emps]

Materialized view

Substitute table scan

```
CREATE MATERIALIZED VIEW  
  emps_by_deptno_gender AS  
SELECT deptno, gender,  
       COUNT(*) AS c, SUM(sal) AS s  
FROM emps  
GROUP BY deptno, gender
```

```
SELECT COUNT(*) AS c  
FROM emps  
WHERE deptno = 10  
AND gender = 'M'
```



Aggregates in LookML

Aggregate table is used even by queries on the orders explore

LookML - orders

```
view: orders {  
  dimension_group: created_at {  
    type: time  
    sql: ${TABLE}.created_at ;;  
  }  
  dimension: product_id {}  
  dimension: customer_id {}  
  measure: revenue {}  
}  
  
view: customers {  
  dimension: customer_id {}  
  dimension: state {}  
}  
  
explore: orders {  
  join: customers {}  
}
```

LookML - orders_by_state_quarter

```
view: orders_by_state_quarter {  
  derived_table: {  
    explore_source: orders {  
      column: created_at_quarter {  
        field: orders.created_at_quarter  
      }  
      column: state {  
        field: customers.state  
      }  
      column: sum_revenue {  
        sql: sum(${orders.revenue});;  
      }  
    }  
  }  
  explore: orders_by_state_quarter {}  
}
```

Advanced features & crazy ideas

- Join elimination
 - If my query uses one table but my aggregate table uses two, is it safe to rewrite?
- Time zones
 - If the aggregate table is rolled up by day, and what if my day ends at a different time than yours?
- Recommend/build aggregate tables based on...
 - Queries over the last month
 - Queries over the last ten minutes
 - Queries needed to satisfy a given dashboard

Agenda

Relational algebra and query optimization

How Looker executes a query (and how Calcite can help)

Aggregate awareness & materialized views

Generate LookML model from SQL

A model out of thin air

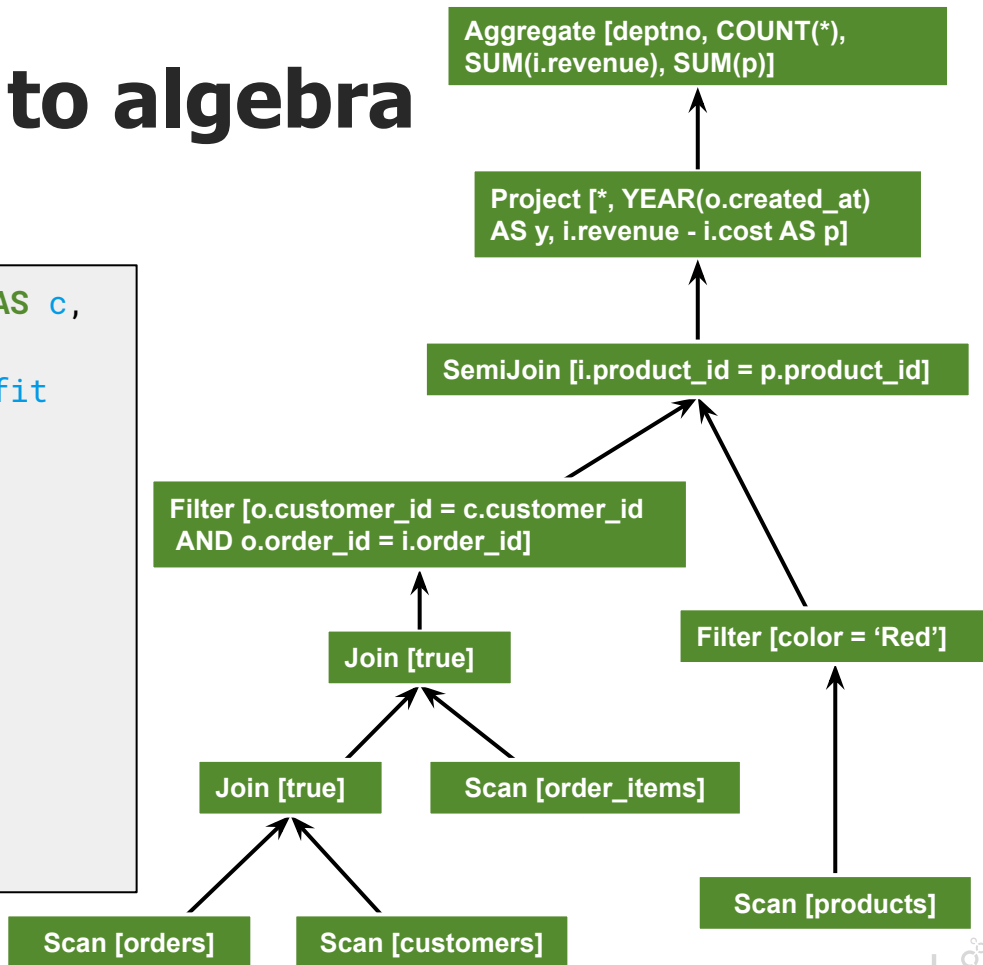
Companies that are just starting to use Looker often have a database and log of queries run against that database.

Parsing the SQL queries in the log, we can recover the dimensions, measures, relationships, and put them into a LookML model.

We call this process “model induction”.

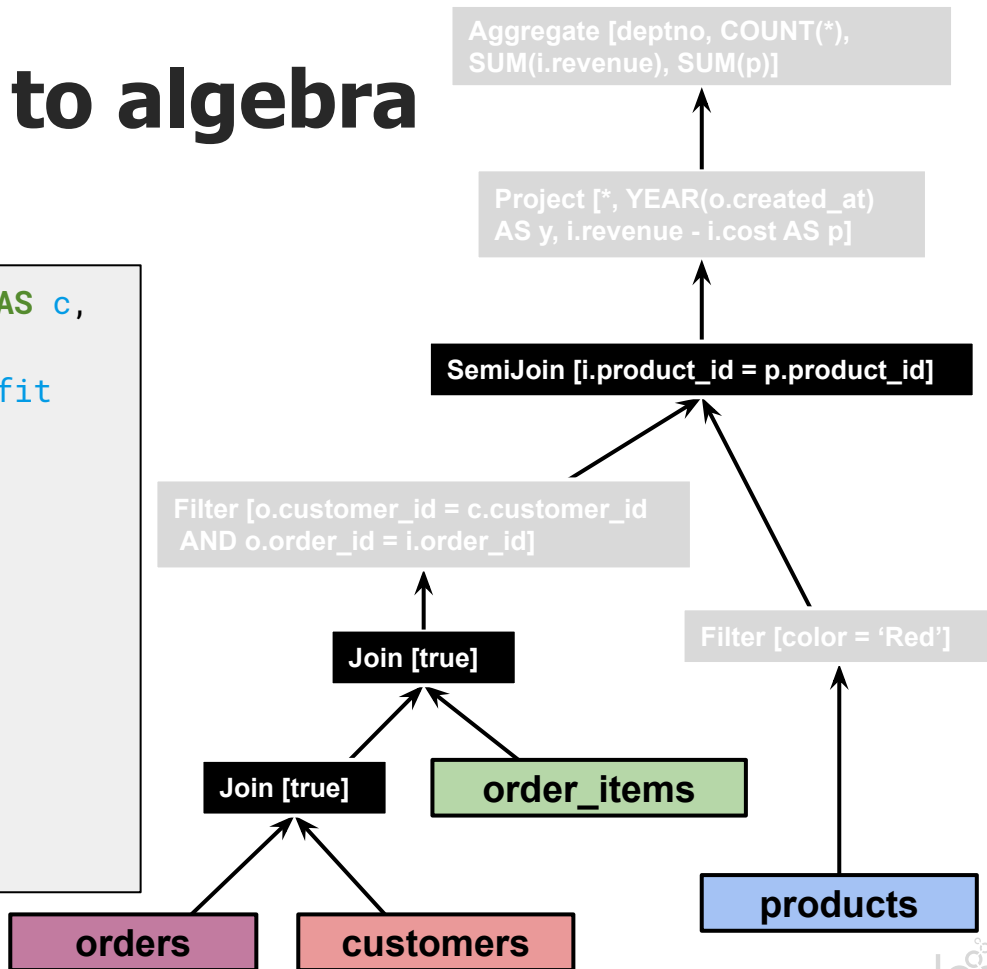
Convert a query to algebra

```
SELECT YEAR(o.created_at), COUNT(*) AS c,  
       SUM(i.revenue) AS sum_revenue,  
       SUM(i.revenue - i.cost) AS sum_profit  
FROM orders AS o,  
     customers AS c,  
     order_items AS i  
WHERE o.customer_id = c.customer_id  
AND o.order_id = i.order_id  
AND i.product_id IN (  
    SELECT product_id  
    FROM products AS p  
    WHERE color = 'Red')  
GROUP BY YEAR(o.created_at)
```



Convert a query to algebra

```
SELECT YEAR(o.created_at), COUNT(*) AS c,  
       SUM(i.revenue) AS sum_revenue,  
       SUM(i.revenue - i.cost) AS sum_profit  
FROM orders AS o,  
     customers AS c,  
     order_items AS i  
WHERE o.customer_id = c.customer_id  
AND o.order_id = i.order_id  
AND i.product_id IN (  
    SELECT product_id  
    FROM products AS p  
    WHERE color = 'Red')  
GROUP BY YEAR(o.created_at)
```




Convert query to join graph, measures

```
SELECT YEAR(o.created_at), COUNT(*) AS c,  
       SUM(i.revenue) AS sum_revenue,  
       SUM(i.revenue - i.cost) AS sum_profit  
FROM orders AS o,  
     customers AS c,  
     order_items AS i  
WHERE o.customer_id = c.customer_id  
AND o.order_id = i.order_id  
AND i.product_id IN (  
    SELECT product_id  
    FROM products AS p  
    WHERE color = 'Red')  
GROUP BY YEAR(o.created_at)
```

measures

COUNT(*)
SUM(i.revenue)
SUM(i.revenue - i.cost)



<show LookML generated from previous query>

Generate LookML model from SQL queries

Query file





```
SELECT d1, d2, m1, m2
FROM customers
JOIN orders
```

```
SELECT d1, d3, m2, m3
FROM products
JOIN orders
```


```
SELECT d2, m4
FROM products
JOIN inventory
```

```
SELECT m5
FROM products
JOIN product_types
```

1. Parse & validate SQL query log; convert to relational algebra



2. Construct join graphs; deduce PK-FK relationships by running cardinality queries



3. Cluster join graphs into stars, each with a central "fact table"

LookML

```
view: customers {}
view: orders {}
view: products {}
view: product_types {}
explore: orders {
  join: customers {}
  join: products {}
  join: product_types {}
}
```

```
view: products {}
view: inventory {}
explore: inventory {
  join: products {}
}
```

4. Generate LookML; one explore per star

< show screenshots or demo of inducer
reading query set and generating model >

< show screenshots or demo of inducer
driven from a single query in SQL Runner >

< a few more possible futures: querying explores via SQL; a JDBC driver as an alternative interface; an internal query language (perhaps SQL, perhaps something else) that allows defining views, explores and queries on top of existing explores >

Thank you! Any questions?

@julianhyde
@ApacheCalcite

<https://calcite.apache.org>
<https://github.com/apache/calcite>

