

# Measures in SQL

Julian Hyde (Google)

2024-05-22

SF Distributed Systems Meetup  
San Francisco, California

## Measures in SQL

Julian Hyde

Google Inc.  
San Francisco, CA, USA  
julianhyde@google.com

John Fremlin

Google Inc.  
New York, NY, USA  
fremlin@google.com

### ABSTRACT

SQL has attained widespread adoption, but Business Intelligence tools still use their own higher level languages based upon a multidimensional paradigm. Composable calculations are what is missing from SQL, and we propose a new kind of column, called a measure, that attaches a calculation to a table. Like regular tables, tables with measures are composable and closed when used in queries.

SQL-with-measures has the power, conciseness and reusability of multidimensional languages but retains SQL semantics. Measure invocations can be expanded in place to simple, clear SQL.

To define the evaluation semantics for measures, we introduce context-sensitive expressions (a way to evaluate multidimensional expressions that is consistent with existing SQL semantics), a concept called evaluation context, and several operations for setting and modifying the evaluation context.

SIGMOD, June 9–15, 2024, Santiago, Chile

# 1. Problem

# Tables are broken!

Tables are unable to provide reusable calculations.

# Problem: Calculate profit margin of orders

| prodName | custName | orderDate  | revenue | cost |
|----------|----------|------------|---------|------|
| Happy    | Alice    | 2023/11/28 | 6       | 4    |
| Acme     | Bob      | 2023/11/27 | 5       | 2    |
| Happy    | Alice    | 2024/11/28 | 7       | 4    |
| Whizz    | Celia    | 2023/11/25 | 3       | 1    |
| Happy    | Bob      | 2022/11/27 | 4       | 1    |

```
SELECT prodName,  
       (SUM(revenue) - SUM(cost))  
       / SUM(revenue) AS profitMargin  
FROM Orders  
WHERE prodName = 'Happy';
```

```
profitMargin  
=====  
0.47
```

# Attempted solution: Create a view

| prodName | custName | orderDate  | revenue | cost |
|----------|----------|------------|---------|------|
| Happy    | Alice    | 2023/11/28 | 6       | 4    |
| Acme     | Bob      | 2023/11/27 | 5       | 2    |
| Happy    | Alice    | 2024/11/28 | 7       | 4    |
| Whizz    | Celia    | 2023/11/25 | 3       | 1    |
| Happy    | Bob      | 2022/11/27 | 4       | 1    |

```
CREATE VIEW SummarizedOrders AS
SELECT prodName, orderDate,
       (SUM(revenue) - SUM(cost))
       / SUM(revenue) AS profitMargin
FROM Orders
GROUP BY prodName, orderDate;
```

```
SELECT prodName,
       (SUM(revenue) - SUM(cost))
       / SUM(revenue) AS profitMargin
FROM Orders
WHERE prodName = 'Happy';
```

```
profitMargin
=====
0.47
```

```
SELECT AVG(profitMargin) AS profitMargin
FROM SummarizedOrders
WHERE prodName = 'Happy';
```

```
profitMargin
=====
0.50
```

## 2. Theory

# Extend the relational model with measures

1. Allow tables to have measures

```
DESCRIBE EnhancedOrders;
```

```
column      type
=====
prodName    STRING
custName    STRING
orderDate   DATE
revenue     INTEGER
cost        INTEGER
profitMargin DOUBLE MEASURE
```

2. Operators for evaluating measures

```
SELECT prodName, profitMargin
FROM EnhancedOrders
GROUP BY prodName;
```

```
prodName profitMargin
=====
Acme      0.60
Happy     0.47
Whizz     0.67
```

3. Syntax to define measures in a query

```
SELECT *,
  (SUM(revenue) - SUM(cost)) / SUM(revenue)
  AS MEASURE profitMargin
FROM Orders
GROUP BY prodName;
```

# Definitions

A **context-sensitive expression** (CSE) is an expression whose value is determined by an evaluation context.

An **evaluation context** is a predicate whose terms are one or more columns from the same table.

- This set of columns is the **dimensionality** of the CSE.

A **measure** is a special kind of column that becomes a CSE when used in a query.

- A measure's dimensionality is the set of non-measure columns in its table.
- The data type of a measure that returns a value of type  $t$  is  $t$  MEASURE, e.g. **INTEGER MEASURE**.

```
SELECT prodName,  
       profitMargin  
FROM EnhancedOrders  
GROUP BY prodName;
```

| prodName | profitMargin |
|----------|--------------|
| Acme     | 0.60         |
| Happy    | 0.50         |
| Whizz    | 0.67         |

```
SELECT (SUM(revenue) - SUM(  
      / SUM(revenue) AS profitMargin  
FROM Orders  
WHERE prodName = 'Acme';  
  
profitMargin  
=====  
0.60
```

profitMargin is a  
measure (and a  
CSE)

Dimensionality is  
{prodName, custName,  
orderDate, revenue,  
cost}

Evaluation context  
for this cell is

prodName = 'Acme'



# AT operator

The **context transformation operator** **AT** modifies the evaluation context.

Syntax:

*expression* **AT** (*contextModifier...*)

*contextModifier* ::=

| **WHERE** *predicate*

| **ALL**

| **ALL** *dimension*

| **SET** *dimension* = [**CURRENT**] *expression*

| **VISIBLE**

```
SELECT prodName,
       profitMargin,
       profitMargin
       AT (SET prodName = 'Happy')
       AS happyMargin,
       profitMargin
       AT (SET custName = 'Bob')
       AS bobMargin
FROM EnhancedOrders
GROUP BY prodName;
```

| prodName | profitMargin | happyMargin | bobMargin |
|----------|--------------|-------------|-----------|
| Acme     | 0.60         | 0.50        | 0.60      |
| Happy    | 0.50         | 0.50        | 0.75      |
| Whizz    | 0.67         | 0.50        | NULL      |

```
SELECT (SUM(revenue) - SUM(cost))
       / SUM(revenue) AS m
FROM Orders
WHERE prodName = 'Whizz'
AND custName = 'Bob';

m
===
NULL
```

Evaluation context for this cell is

prodName = 'Whizz'  
AND custName = 'Bob'

# 3. Consequences

# Grain-locking

What is the average age of the customer who would ordered each product?

When we use an aggregate function in a join query, it will 'double count' if the join duplicates rows.

This is generally not we want for measures – except if we want a weighted average – but is difficult to avoid in SQL.

Measures are locked to the grain of the table that defined them.

| prodName | custName | orderDate  | revenue | cost |
|----------|----------|------------|---------|------|
| Happy    | Alice    | 2023/11/28 | 6       | 4    |
| Acme     | Bob      | 2023/11/27 | 5       | 2    |
| Happy    | Alice    | 2024/11/28 | 7       | 4    |
| Whizz    | Celia    | 2023/11/25 | 3       | 1    |
| Happy    | Bob      | 2022/11/27 | 4       | 1    |

| custName | custAge |
|----------|---------|
| Alice    | 23      |
| Bob      | 41      |
| Celia    | 17      |

```
WITH EnhancedCustomers AS (  
  SELECT *,  
    AVG(custAge) AS MEASURE avgAge  
  FROM Customers)  
SELECT o.prodName,  
  AVG(c.custAge) AS weightedAvgAge,  
  c.avgAge AS avgAge  
FROM Orders AS o  
JOIN EnhancedCustomers AS c USING (custName)  
GROUP BY o.prodName;
```

| prodName | weightedAvgAge | avgAge |
|----------|----------------|--------|
| =====    | =====          | =====  |
| Acme     | 41             | 41     |
| Happy    | 29             | 32     |
| Whizz    | 17             | 17     |

Alice (age 23)  
has two orders;  
Bob (age 41) has  
one order.

# Composition & closure

Just as tables are closed under queries, so tables-with-measures are closed under queries-with-measures

Measures can reference measures

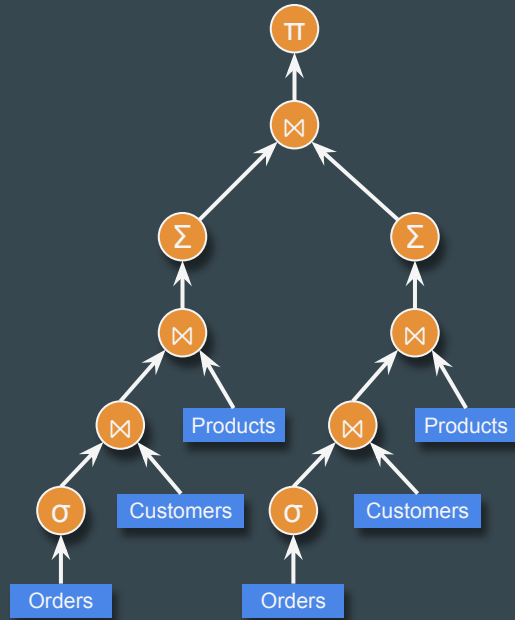
Complex analytical calculations without touching the **FROM** clause

Evaluation contexts can be nested

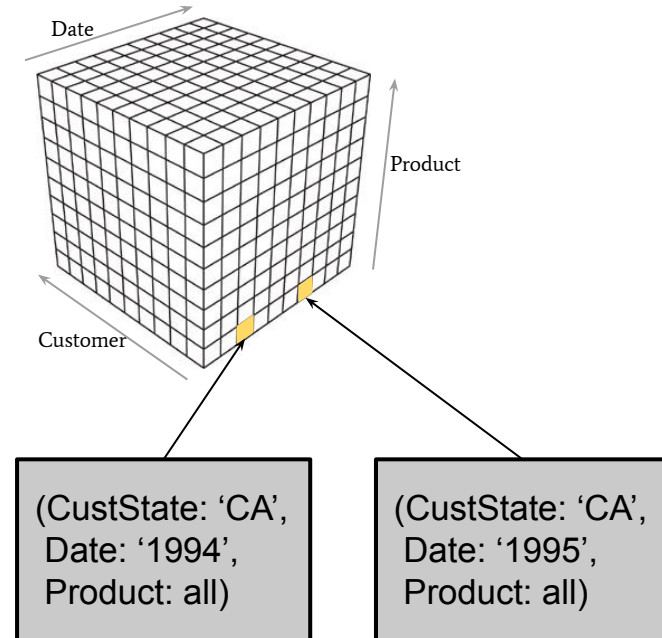
```
SELECT *,
  SUM(cost) AS MEASURE sumCost,
  SUM(revenue) AS MEASURE sumRevenue,
  (sumRevenue - sumCost) / sumRevenue
    AS MEASURE profitMargin,
  sumRevenue
    - sumRevenue AT (SET YEAR(orderDate)
      = CURRENT YEAR(orderDate) - 1)
    AS MEASURE revenueGrowthYoY,
  ARRAY_AGG(productId
    ORDER BY sumRevenue DESC LIMIT 5)
    AT (ALL productId)
    AS MEASURE top5Products,
  ARRAY_AGG(customerId
    ORDER BY sumRevenue DESC LIMIT 3)
    AT (ALL customerId
      SET productId MEMBER OF top5Products
        AT (SET YEAR(orderDate)
          = CURRENT YEAR(orderDate) - 1))
    AS MEASURE top3CustomersOfTop5Products
FROM Orders;
```

# Bottom-up vs Top-down query


Relational algebra (bottom-up)



Multidimensional (top-down)



# Represent a Business Intelligence model as a SQL view



- SQL planner handles view expansion
- Grain locking makes it safe to use a star schema
- Users can define new models simply by writing queries

```
CREATE VIEW OrdersCube AS
SELECT *
FROM (
  SELECT o.orderDate AS `order.date`,
         o.revenue AS `order.revenue`,
         SUM(o.revenue) AS MEASURE `order.sum_revenue`
  FROM Orders) AS o
LEFT JOIN (
  SELECT c.custName AS `customer.name`,
         c.state AS `customer.state`,
         c.custAge AS `customer.age`,
         AVG(c.custAge) AS MEASURE `customer.avg_age`
  FROM Customers) AS c
ON o.custName = c.custName
LEFT JOIN (
  SELECT p.prodName AS `product.name`,
         p.color AS `product.color`,
         AVG(p.weight) AS MEASURE `product.avg_weight`
  FROM Products) AS p
ON o.prodName = p.prodName;
```

```
SELECT `customer.state`, `product.avg_weight`
FROM OrdersCube
GROUP BY `customer.state`;
```

# Implementing measures & CSEs as SQL rewrites

simple



complex

| Complexity                              | Query  | Expanded query  |
|---|--|---|
| Simple measure can be inlined           | <pre>SELECT prodName, avgRevenue FROM OrdersCube GROUP BY prodName</pre>   | <pre>SELECT prodName, AVG(revenue) FROM orders GROUP BY prodName</pre>  |
| Join requires grain-locking             | <pre>SELECT prodName, avgAge FROM OrdersCube GROUP BY prodName</pre>   | <pre>SELECT o.prodName, AVG(c.custAge PER c.custName) FROM orders JOIN customers GROUP BY prodName → (something with GROUPING SETS)</pre>             |
| Period-over- period                     | <pre>SELECT prodName, avgAge -     avgAge AT (SET year =         CURRENT year - 1) FROM OrdersCube GROUP BY prodName</pre>                           | (something with window aggregates)  |
| Scalar subquery can accomplish anything | <pre>SELECT prodName, prodColor     avgAge AT (ALL custState         SET year = CURRENT year - 1) FROM OrdersCube GROUP BY prodName, prodColor</pre> | <pre>SELECT prodName, prodColor,     (SELECT ... FROM orders         WHERE &lt;evaluation context&gt;) FROM orders GROUP BY prodName, prodColor</pre> |


# Summary

Top-down evaluation makes queries concise

Measures make calculations reusable

Measures don't break SQL





Thank you!  
Any questions?

@julianhyde

@JohnFremlin

@ApacheCalcite

<https://calcite.apache.org>

<https://issues.apache.org/jira/browse/CALCITE-4496>

<http://www.hydromatic.net/hyde2024measures.pdf>