

Measures in SQL



arxiv.org/abs/2406.00251

- Composable calculations are what is missing from SQL. We propose a new kind of column, called a measure, that attaches a calculation to a table.
- Like regular tables, tables with measures are composable and closed when used in queries.
- SQL-with-measures has the power,
 - conciseness
 - and reusabilityof multidimensional languages – but retains SQL semantics.
- Measure invocations can be expanded in place to simple, clear SQL.



Extending relational model

Measure calculations are holographic values defined as columns with **AS MEASURE**.

They are context sensitive expressions.

The **context transformation operator AT** modifies the evaluation context.

Syntax:

expression AT (contextModifier...)
contextModifier ::= WHERE predicate

- | ALL
- | ALL *dimension*
- | SET *dimension* = [CURRENT] *expression*
- | VISIBLE

prodName	custName	orderDate	revenue	cost
Happy	Alice	2023/11/28	6	4
Acme	Bob	2023/11/27	5	2
Happy	Alice	2024/11/28	7	4
Whizz	Celia	2023/11/25	3	1
Happy	Bob	2022/11/27	4	1

Composing measures

```
SELECT *,
SUM(cost) AS MEASURE sumCost,
SUM(revenue) AS MEASURE sumRevenue,
(sumRevenue - sumCost) / sumRevenue
AS MEASURE profitMargin,
sumRevenue
- sumRevenue AT (SET YEAR(orderDate)
= CURRENT YEAR(orderDate) - 1)
AS MEASURE revenueGrowthYoY,
ARRAY_AGG(prodName
ORDER BY sumRevenue DESC LIMIT 5)
AT (ALL prodName)
AS MEASURE top5Products,
ARRAY_AGG(custName
ORDER BY sumRevenue DESC LIMIT 3)
AT (ALL custName
SET productId MEMBER OF top5Products
AT (SET YEAR(orderDate) - 1))
= CURRENT YEAR(orderDate) - 1))
AS MEASURE top3CustomersOfTop5Products
FROM Orders;
```

Grain locking

A measure is locked to the grain of its defining table.

Joining another table does not introduce double-counting.

Measure invocations can still be rewritten into SQL without measures via a semi-join (WHERE EXISTS).

custName	custAge
Alice	23
Bob	41
Celia	17

prodName	custName	orderDate	revenue	cost	custAge
Happy	Alice	2023/11/28	6	4	23
Acme	Bob	2023/11/27	5	2	41
Happy	Alice	2024/11/28	7	4	23
Whizz	Celia	2023/11/25	3	1	17
Happy	Bob	2022/11/27	4	1	41

Alice appears twice affecting the average of custAge after the JOIN

JOIN query with measure

```
WITH EnhancedCustomers AS (
SELECT *,
AVG(custAge) AS MEASURE avgAge
FROM Customers)
SELECT o.prodName,
AVG(c.custAge) AS weightedAvgAge,
c.avgAge AS avgAge
FROM Orders AS o
JOIN EnhancedCustomers AS c USING
(custName)
GROUP BY o.prodName;
```

Output: avgAge measure keeps grain

prodName	weightedAvgAge	avgAge
Acme	41	41
Happy	29 (Alice counted twice)	32
Whizz	17	17

Measure rewritten to semi-JOIN

```
SELECT
o.prodName,
AVG(c.custAge) AS weightedAvgAge,
(
SELECT AVG(c2.custAge)
FROM Customers AS c2
WHERE
EXISTS(
SELECT TRUE FROM Orders AS o2
WHERE c2.custName = o2.custName join condition
AND o.prodName = o2.prodName grouping key filter context
)
) AS avgAge
FROM Orders AS o
JOIN Customers AS c
USING (custName)
GROUP BY o.prodName
```

Rewriting

Implementing measures & context sensitive expressions as SQL rewrites

	Complexity	Query	Expanded query
simple	Simple measure can be inlined	SELECT prodName, avgRevenue FROM OrdersCube GROUP BY prodName	SELECT prodName, AVG(revenue) FROM orders GROUP BY prodName
	Join requires grain-locking	SELECT prodName, avgAge FROM OrdersCube GROUP BY prodName	SELECT o.prodName, AVG(ANY_VALUE(c.custAge) GROUP BY c.custName) FROM orders JOIN customers GROUP BY prodName → (something with GROUPING SETS)
	Period-over- period	SELECT prodName, avgAge - avgAge AT (SET year = CURRENT year - 1) FROM OrdersCube GROUP BY prodName	(something with window aggregates)
complex	Scalar subquery can accomplish anything	SELECT prodName, prodColor avgAge AT (ALL custState SET year = CURRENT year - 1) FROM OrdersCube GROUP BY prodName, prodColor	SELECT prodName, prodColor, (SELECT ... FROM orders WHERE <evaluation context>) FROM orders GROUP BY prodName, prodColor

Let's connect!

Topic ideas to chat about (more SQL innovations in the pipes!)

- Formal semantics for filter context transformations
- Using measures with LLM assistants
- Operators for managing grain
- Implementation strategies (efficient grain locking)
- Changing temporal grain for forecasts and timeseries
- Adding measures to nested structures typical in log files
- Sequential processing and measures



Julian Hyde
@julianhyde @ApacheCalcite
julianhyde@google.com



John Fremlin
@johnfremlin
fremlin@google.com