

# More than Query


Future Directions of Query Languages,  
from SQL to Morel



**Julian Hyde** (ex-Google; founder of Apache Calcite & Morel)

April 23rd, 2025 4:15 pm

Oakland



# Solving Wordle

What is the best word to guess next?

Guess: WHARF

W H A R F  
W H A R F

CAIRN CARRY COBRA +57

BOARD DIARY UNARY +19

+ 22 groups

Guess: MORAL

M O R A L  
M O R A L

DRYAD RAJAH URBAN +4

BURAN HYRAX YDRAD +10

+ 47 groups

For each possible guess,

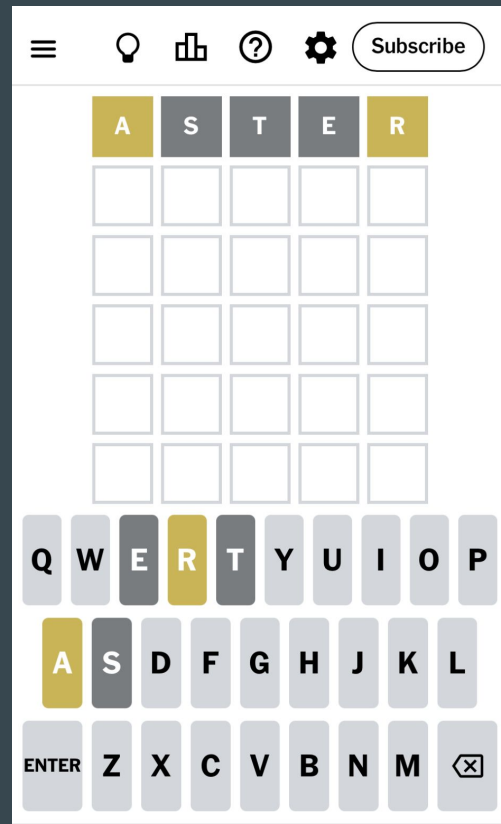
For each possible answer,

Compute the mask for that guess for that answer

Group answers that have the same mask

Count the groups

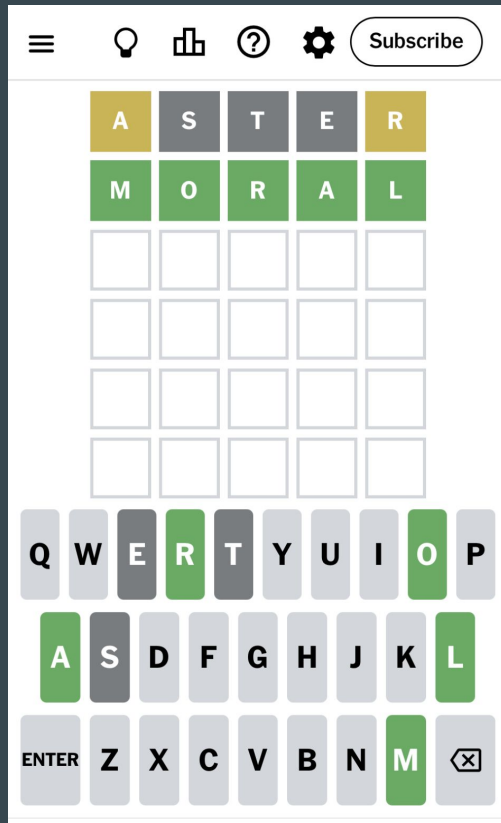
Choose the guess that has the most groups



# Solving Wordle

morel version 0.5.0 (java version "21.0.6", JRE null (build 21.0.6+8-LTS-188), JLine terminal, xterm-256color)

```
- use "wordle.sml";
[opening wordle.sml]
val words =
  ["aahed", "aalii", "aargh", "aarti", "abaca", "abaci", "aback", "abacs", "abaft",
   "abaka", "abamp", "aband", "abase", "abash", "abask", "abate", "abaya", "abbas",
   "abbed", "abbes", ...] : string list
val mask = fn : string * string -> int
val maskToString = fn : int -> string
val maskCount = fn : string * string list -> int
val bestGuesses = fn : string list -> {maskCount:int, w:string} list
val remaining = fn : string list * (string * string) list -> string list
-
- bestGuesses (remaining (words, [("aster", "ybbby")]));
val it =
  [{maskCount=54, w="marid"}, {maskCount=54, w="moria"}, {maskCount=51, w="daric"},
   {maskCount=51, w="moira"}, {maskCount=50, w="caird"}, {maskCount=50, w="coria"},
   {maskCount=49, w="carbo"}, {maskCount=49, w="cardi"}, {maskCount=49, w="coram"},
   {maskCount=49, w="darcy"}, {maskCount=49, w="moral"}, {maskCount=48, w="maria"},
   {maskCount=47, w="baric"}, {maskCount=47, w="borna"}, {maskCount=47, w="cairn"},
   {maskCount=47, w="cardy"}, {maskCount=47, w="carny"}, {maskCount=47, w="caron"},
   {maskCount=47, w="cobra"}, {maskCount=47, w="coral"}, ...]
: {maskCount:int, w:string} list
```



# Solving Wordle

```
fun mask (guess, answer) =  
  Let  
    fun mask2 (m, i, [], answer) = m  
      | mask2 (m, i, letter :: rest, answer) =  
        mask2 ((m * 3  
          + (if sub(answer, i) = letter  
            then 2  
            else if isSubstring(str letter) answer  
              then 1  
              else 0)), i + 1, rest, answer)  
    in  
      mask2 (0, 0, explode guess, answer)  
    end;  
fun maskToString m =  
  let  
    fun maskToString2 (m, s, 0) = s  
      | maskToString2 (m, s, k) =  
        maskToString2 (m div 3,  
          List.nth(["b", "y", "g"], m mod 3) ^ s,  
            k - 1)  
    in  
      maskToString2 (m, "", 5)  
    end;
```

```
val words =  
  from w in file.wordle.words yield w.word;  
fun maskCount (guess, remainingWords) =  
  from w in remainingWords  
    group m = mask (guess, w) compute c = count  
    compute count;  
fun bestGuesses words =  
  from w in words,  
    maskCount = maskCount (w, words)  
    order maskCount desc;  
fun remaining (words, []) = words  
  | remaining (words, (guess, m) :: rest) =  
    from w in (remaining (words, rest))  
      where maskToString (mask (guess, w)) = m;
```



Subscribe

A

S

T

E

R

P

L

ENTER

Z

X

C

V

B

N

M



# Agenda

What's the difference between a query language and a programming language?

What's wrong with SQL?

What things do other query-like languages do better?

Morel, a functional query language

# Algebra – a collection of values, operators, and laws

## On numbers

$$a + (b + c) \rightarrow (a + b) + c$$

$$(a * b) + (a * c) \rightarrow a * (b + c)$$

## On relations

$$P \cup (Q \cup R) \rightarrow (P \cup Q) \cup R$$

$$(P \bowtie Q) \cup (P \bowtie R) \rightarrow P \bowtie (Q \cup R)$$

# Functional programming – values, types, operators

```
1 + 2;  
> val it = 3 : int  
  
"Hello, " ^ "world!";  
> val it = "Hello, world!" : string  
  
val integers = [1, 2, 3, 4, 5, 6, 7, 8];  
> val integers = [1,2,3,4,5,6,7,8] : int list  
  
fun filter f [] = []  
  | filter f (first :: rest) =  
    if (f first)  
    then first :: (filter f rest)  
    else filter f rest;  
val filter = fn : ('a -> bool) -> 'a list -> 'a list  
  
filter (fn i => i mod 2 = 0) integers;  
> val it = [2,4,6,8] : int list
```

# Relational algebra in a functional programming language

## Relational algebra

$\cup$  union  
 $\setminus$  minus  
 $\cap$  intersect  
 $\sigma$  filter  
 $\Pi$  project  
 $\bowtie$  join

## Relational operators as functions

```
val union = fn : 'a list * 'a list -> 'a list
val except = fn : 'a list * 'a list -> 'a list
val intersect = fn : 'a list * 'a list -> 'a list
val filter = fn : ('a -> bool) -> 'a list -> 'a list
val map = fn : ('a -> 'b) -> 'a list -> 'b list
val join = fn
  : 'a list * 'b list * ('a * 'b -> bool)
  -> ('a * 'b) list
```



**Is SQL a functional programming language?**

## Query

- Built-in operators
- Relation is the only bulk type
- Algebraic rewrite
- Not Turing complete
- Distributed engine

## Query + functional

- Immutable data
- Type inference
- Inlining

## Functional

- User-defined operators
- User-defined types
- Functions as values
- Local optimizations
- Turing complete
- Local execution

## Algebraic

- User-defined operators & types
- Algebraic rewrite
- Turing complete
- Choose your own engine

The image features a word cloud centered around the acronym 'SQL'. The background is a deep space scene with a prominent spiral galaxy in shades of blue and white, set against a black field of distant stars. A semi-transparent blue circle is positioned behind the central text. The words are rendered in a green, monospace-style font. The central word 'SQL' is larger and in a light orange color. Other words are arranged in a circular pattern around it, with varying sizes and orientations. The words include: arithmetic, optimization, constraint programming, linear algebra, vector, statistics, analytics, reporting, search, transaction processing, general programming, engineering, data, deductive, graph, geospatial, and sat.

arithmetic  
optimization  
constraint programming  
linear algebra  
vector  
statistics  
analytics  
reporting  
search  
transaction processing  
general programming  
engineering  
data  
deductive  
graph  
geospatial  
SQL  
sat

# Future of SQL

I believe:

- SQL continues to grow/sprawl
- ISO standard declines in importance
- Databases speak more than one language
- Databases are not the only systems that speak SQL


SQL as the only API




Federation / virtualization



SQL as machine code



Polyglot systems



# Alternatives to SQL

## SQL

```
SELECT name, sal + comm AS pay
FROM emps
WHERE deptno = 10
```

## MongoDB (JavaScript)

```
db.emps.find(
  { deptno: 10 },
  { name: 1, pay: { $add: ["$sal", "$comm"] }, _id: 0 }
)
```

## Datomic (Clojure)

```
(d/q '[:find ?name ?pay
      :where
      [?e :emps/deptno 10]
      [?e :emps/name ?name]
      [?e :emps/sal ?sal]
      [?e :emps/comm ?comm]
      [(+ ?sal ?comm) ?pay]]
  db)
```

## Spark (Scala)

```
df.filter($"deptno" === 10)
  .select($"name", ($"sal" + $"comm").alias("pay"))
```

## Pandas (Python)

```
df[(df['deptno'] == 10)][['name', 'sal', 'comm']] \
  .assign(pay=lambda x: x['sal'] + x['comm'])[['name', 'pay']]
```



Morel

```
from e in db.emps
  where e.deptno = 2
  yield {e.name, pay = e.sal + e.comm}
```

XQuery

```
for $emp in collection("emps")/employee
where $emp/deptno = 10
return
  <result>
    <name>{$emp/name/text()}</name>
    <pay>{$emp/sal/text() + $emp/comm/text()}</pay>
  </result>
```

LINQ (C#)

```
var result = from emp in emps
              where emp.deptno == 10
              select new {
                  name = emp.name,
                  pay = emp.sal + emp.comm
              };
```

# SQL

```
-- SQL
SELECT item, COUNT(*) AS c,
       SUM(sales) AS total
FROM ProduceSales
WHERE item != 'bananas'
      AND category IN ('fruit', 'nut')
GROUP BY item
ORDER BY item DESC;
```

item	c	total
=====	=====	=====
apples	2	9

## GoogleSQL pipe syntax

```
-- GoogleSQL pipe syntax
FROM ProduceSales
|> WHERE item != 'bananas'
      AND category IN ('fruit', 'nut')
|> AGGREGATE COUNT(*) AS c, SUM(sales) AS total
      GROUP BY item
|> ORDER BY item DESC;
```



# SQL

```
-- SQL
SELECT item, COUNT(*) AS c,
       SUM(sales) AS total
FROM ProduceSales
WHERE item != 'bananas'
      AND category IN ('fruit', 'nut')
GROUP BY item
ORDER BY item DESC;
```

item	c	total
=====	=====	=====
apples	2	9

# PRQL

```
# PRQL
from produceSales
filter item != "bananas"
filter category in ["fruit", "nut"]
group item (
  aggregate {
    c = count this,
    total = sum sales
  }
)
sort -item
```

# SQL

```
-- SQL
SELECT item, COUNT(*) AS c,
       SUM(sales) AS total
FROM ProduceSales
WHERE item != 'bananas'
      AND category IN ('fruit', 'nut')
GROUP BY item
ORDER BY item DESC;
```

item	c	total
=====	=====	=====
apples	2	9



```
(* Morel *)
from p in produceSales
  where p.item != "bananas"
    andalso p.category elem ["fruit", "nut"]
  group p.item compute c = count,
    total = sum of p.sales
  order item desc
```

# Morel query syntax

```
exp →  
    from scan [, scan... ] [ step... ]  
    [ terminalStep ]  
    | exists scan [, scan... ] [ step... ]  
    | forall scan [, scan... ] [ step... ]  
    require condition  
    | (* many other kinds of expression *)
```

```
scan →  
    pat in collection [ on condition ]  
    | pat = exp [ on condition ]  
    | var
```

```
key → [ id = ] exp
```

```
agg → [ id = ] aggFn of exp
```

```
orderItem → exp [ desc ]
```

```
step →  
    join scan [, scan... ]  
    | where condition  
    | yield exp  
    | yieldmany collection  
    | distinct  
    | group key [, key... ]  
    [ compute agg [, agg... ] ]  
    | order item [, item... ]  
    | skip number  
    | take number  
    | through pat in exp  
    | union collection [, collection... ]  
    | intersect collection [, collection... ]  
    | except collection [, collection... ]
```

```
terminalStep →  
    into exp  
    | compute agg [, agg... ]
```

# The billion-dollar mistake

```
SELECT e.ename, e.sal
FROM emp AS e
WHERE e.deptno = 10
AND e.sal >
    (SELECT MAX(e2.sal)
     FROM emp AS e2
     WHERE e2.deptno = 20
     AND e2.job = 'PROGRAMMER')
```

# The billion-dollar mistake

Why do we need NULL?

- Empty strings
- Missing references
- Missing values
- N/A
- Unknown boolean
- Aggregation over the empty set

We don't need NULL!

- Empty strings
- Option type
- Option type
- Option type
- Unnecessary
- Aggregate functions that return a default value, or an option

```
SELECT e.ename, e.sal
FROM emps AS e
WHERE e.deptno = 10
AND e.sal >
    (SELECT MAX(e2.sal)
     FROM emps AS e2
     WHERE e2.deptno = 20
     AND e2.job = 'PROGRAMMER')
```

```
from e in scott.emp
  where e.deptno = 10
  andalso
    (forall e2 in scott.emp
      where e2.deptno = 20
      andalso e2.job = "PROGRAMMER"
      require e.sal > e2.sal)
  yield e.ename
```

# Types and collections



## Atomic types

- **Primitive types** – e.g. `bool`, `char`, `string`, `int`, `real`
- **Abstract types** – e.g. `instant`, `interval`, `point`

## Collection types

- **Set** – no ordering, no duplicates
- **List** – ordering, duplicates allowed
- **Bag** – no ordering, duplicates allowed

A relation is a *set of records* (academia), a *bag of records* (industry)

## Functions and polymorphism

- **Type variables** – enable polymorphic types
- **Function types** – e.g. `int → bool`,  
`'a → 'b → 'a list → 'b list`

## Algebraic types

- **Record** – each value has several named fields
- **Sum type** – each value is one of several named variants

```
datatype personnel_id =  
    EMPLOYEE of int  
  | CONTRACTOR of {ssid: string, agency: string};  
  
type member = {name: string, deptno: int, id: personnel_id};  
  
val members = [  
    {name = "Smith", deptno = 10, id = EMPLOYEE 100},  
    {name = "Jones", deptno = 20,  
     id = CONTRACTOR {ssid = "xxx-xx-xxxx", agency = "Cheap & cheerful"}}];  
  
val departments = scott.depts;  
  
val primes = [2, 3, 5, 7, 11];  
  
val bands = [ ["john", "paul", "george", "ringo"], ["simon", "garfunkel"] ];
```

# Lists, bags, sets in Morel

Morel has **list** and **bag** types, but no **set** type. (**list** and **bag** have different operators; **set** is just **bag** plus constraints.)

Rules for a **from** expression:

- The output of a join or set operation is a **list** if all inputs are **list** values, otherwise a **set**;
- **order** keyword converts **bag** (or **list**) to **list**;
- **unordered** converts **list** (or **bag**) to **bag**;
- other operators (**where**, **group**, **distinct**, **yield**, **take**, **skip**) preserve collection type.

```
from i in [1, 2],  
    j in ["a", "b"],  
    k in [3, 4, 5, 6]  
    where i + k < 6;  
{i: int, j: string, k: int} list
```

```
from dept in [10, 30],  
    e in scott.emps  
    where e.deptno = dept  
    yield e.ename;  
string bag
```

```
from dept in [10, 30],  
    e in scott.emps  
    where e.deptno = dept  
    order e.sal desc  
    take 3  
    yield {e.deptno, e.ename};  
{deptno: int, ename: string} list
```

# Modifying data

# SQL DML

*-- Delete employees who earn more than 1,000.*

```
DELETE FROM scott.emps  
WHERE sal > 1000;
```

*-- Add one employee.*

```
INSERT INTO scott.emps (empno, deptno, ename, job, sal)  
VALUES (100, 20, 'HYDE', 'ANALYST', 1150);
```

*-- Double the salary of all managers.*

```
UPDATE scott.emps  
SET sal = sal * 2  
WHERE job = 'MANAGER';
```

*-- Commit.*

```
COMMIT;
```

# Morel DML (first attempt)

```
(* Delete employees who earn more than 1,000. *)
```

```
delete e in scott.emps  
  where e.sal > 1000;
```

```
(* Add one employee. *)
```

```
insert scott.emps  
  [{empno = 100, deptno = 20, ename = "HYDE",  
    job = "ANALYST", sal = 1150}];
```

```
(* Double the salary of all managers. *)
```

```
update e in scott.emps  
  where e.job = "MANAGER"  
  assign (e, {e with sal = e.sal * 2});
```

```
(* Commit. *)
```

```
commit;
```

# Morel DML

```
(* Delete employees who earn more than 1,000. *)
val emps2 =
  from e in scott.emps
    where not (e.sal > 1000);

(* Add one employee. *)
val emps3 = emps2 union
  [{empno = 100, deptno = 20, ename = "HYDE", job = "ANALYST",
    sal = 1150}];

(* Double the salary of all managers. *)
val emps4 =
  from e in emps3
    yield if e.job = "MANAGER"
      then {e with sal = e.sal * 2}
      else e;

(* Commit. *)
commit {scott with emps = emps4};
```



# Incremental computation

```
(* New and removed employees. *)  
val empsAdded = emps4 except scott.emps;  
val empsRemoved = scott.emps except emps4;  
  
(* Compute the updated summary table. *)  
val summary2 =  
  from s in scott.summary  
  union  
    (from e in empsAdded  
     yield {e.deptno, c = 1, sum_sal = e.sal})  
  union  
    (from e in empsRemoved  
     yield {e.deptno, c = ~1, sum_sal = ~sum_sal})  
  group s.deptno compute c = sum of c, sum_sal = sum of sum_sal  
  where c != 0);  
  
(* Commit. *)  
commit {scott with summary = summary2};
```

# Optimizing data-intensive programs

Query optimization (e.g. change join order, push down filters, parallelize)

Program optimization (e.g. remove unused variables, unwind loops, convert tail-recursion to loop)

Efficiently compute deltas

Recommend materialized views

Add temporary tables for shared intermediate results

Remove unused temporary tables

Multi-query optimization  
[[CALCITE-6188](#)]

# Deductive query, Graph query, and Optimization problems

# Parents (a base relation)

```
-- SQL
SELECT *
FROM parents;

parent  child
=====
earendil elrond
elrond   arwen
elrond   elladan
elrond   elrohir

SELECT *
FROM parents
WHERE parent = 'elrond';

parent child
=====
elrond arwen
elrond elladan
elrond elrohir
```

```
/* Datalog */
is_parent(earendil, elrond).
is_parent(elrond, arwen).
is_parent(elrond, elladan).
is_parent(elrond, elrohir).

answer(X) :- is_parent(elrond, X).
X = arwen
X = elladan
X = elrohir
```

```
(* Morel *)
from (parent, child) in parents
  where parent = "elrond";
[("elrond", "arwen"),
 ("elrond", "elladan"),
 ("elrond", "elrohir")]
```

# Ancestors (a recursively-defined relation)

```
CREATE VIEW ancestors AS
WITH RECURSIVE a AS (
  SELECT parent AS ancestor,
         child AS descendant
  FROM parents
  UNION ALL
  SELECT a.ancestor, p.child
  FROM parents AS p
  JOIN a ON a.descendant = p.parent)
SELECT * FROM a;
```

```
SELECT *
FROM ancestors
WHERE descendant = 'arwen';
```

```
ancestor descendant
=====
earendil arwen
elrond arwen
```

StackOverflowError  
is not easy to fix. It's  
hard to write a  
recursive function that  
iterates until a set  
reaches a fixed point.

```
is_ancestor(X, Y) :- is_parent(X, Y).
is_ancestor(X, Y) :- is_parent(X, Z),
is_ancestor(Z, Y).
```

```
answer(X) :- is_ancestor(X, arwen).
X = elrond
X = earendil
```

```
fun ancestors () =
  (from (x, y) in parents)
  union
  (from (x, y) in parents,
   (y2, z) in ancestors ()
   where y = y2
   yield (x, z));
```

```
from (ancestor, descendant) in ancestors ()
  where descendant = "arwen";
Uncaught exception: StackOverflowError
```

# Two ways to define a relation

```
(* Morel "forwards" relation *)

(* Relation defined using algebra. *)
fun clerks () =
  from e in emps
    where e.job = "CLERK";

(* Query uses regular iteration. *)
from e in clerks,
  d in depts
  where d.deptno = e.deptno
    andalso d.loc = "DALLAS"
  yield e.name;
["SMITH", "ADAMS"] : string list;
```

```
(* Morel "backwards" relation *)

(* Relation defined using a predicate. *)
fun isClerk e =
  e.job = "CLERK";

(* Query uses a mixture of constrained
   and regular iteration. *)
from e,
  d in depts
  where isClerk(e)
    andalso d.deptno = e.deptno
    andalso d.loc = "DALLAS"
  yield e.name;
["SMITH", "ADAMS"] : string list;
```

Unbounded variable **e** iterates over all values that could yield a result.

# Recursively-defined predicate relation

```
(* More1 *)  
  
fun isAncestor (x, z) =  
  (x, z) elem parents  
  orelse (exists y  
    where isAncestor(x, y)  
      andalso (y, z) elem parents);  
  
from a  
  where isAncestor(a, "arwen");  
["earendil", "elrond"] : string list  
  
from d  
  where isAncestor("earendil", d);  
["elrond", "arwen", "elladan", "elrohir"] : string list
```

# Recipe optimization

```
from b, c
  where b >= 0 andalso b <= 100      (*) number of banana cakes
  andalso c >= 0 andalso c <= 100    (*) number of chocolate cakes
  andalso 50 * b + 200 * c <= 40000  (*) flour
  andalso 2 * b <= 60                (*) bananas
  andalso 75 * b + 150 * c <= 20000  (*) sugar
  andalso 100 * b + 150 * c <= 5000  (*) butter
  andalso 75 * c <= 5000             (*) cocoa
  yield {b, c, profit = 400 * b + 450 * c}
  order profit desc
  take 3;
val it =
  [{b=29,c=14,profit=17900},
   {b=30,c=13,profit=17850},
   {b=26,c=16,profit=17600}] : {b:int, c:int, profit:int} list
```

## Available ingredients

- 40kg flour
- 60 bananas
- 20kg sugar
- 5kg butter
- 5kg cocoa

## Banana cake \$4

- 50g flour
- 2 bananas
- 75g sugar
- 100g butter



## Chocolate cake \$4.50

- 200g flour
- 150g sugar
- 150g sugar
- 75g cocoa





**If only... using a database could be more like  
using a programming language**

# Many nice things about programming languages

If it compiles, it probably works

Refactoring

Autocompletion

Git

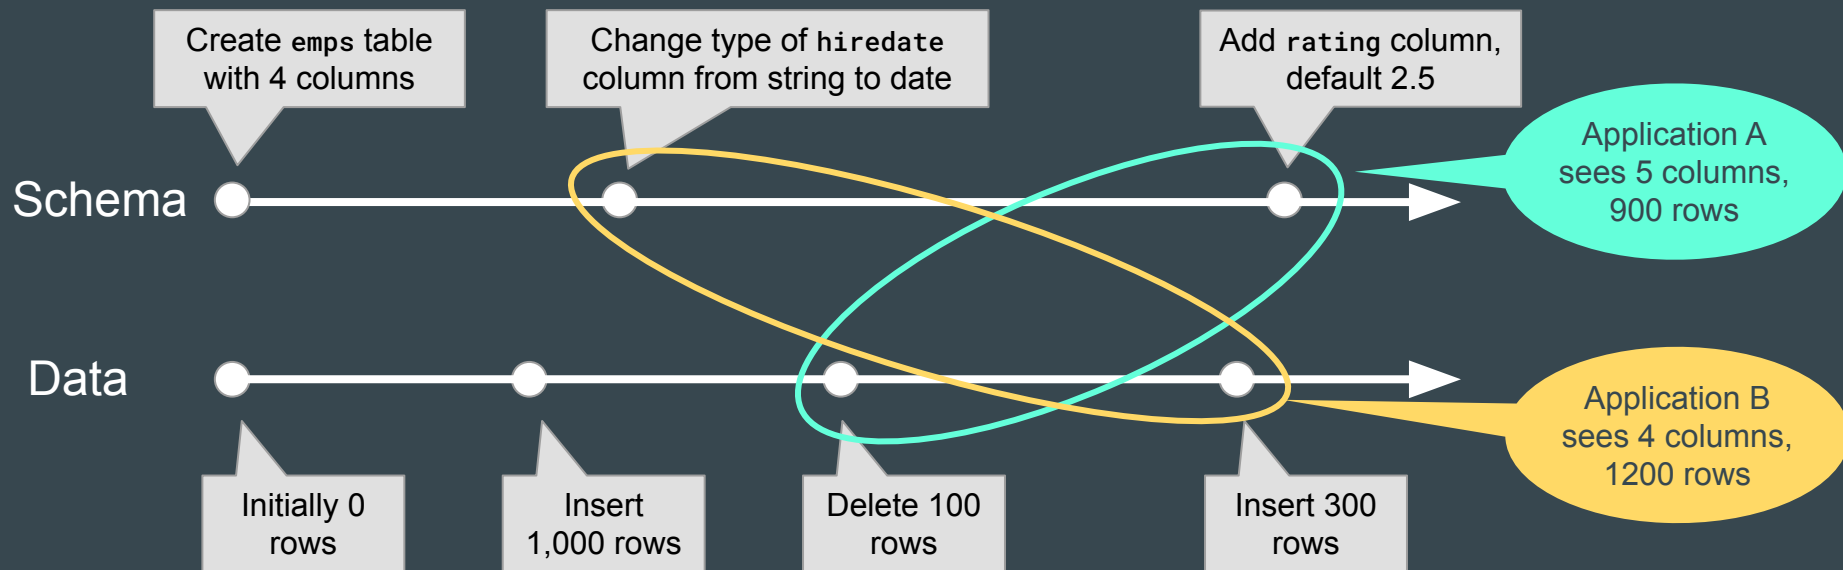
Documentation in the code

Unit tests in the same language

Modules

Abstraction

# Types and data evolve independently



# Database as a module (shims for schema evolution)


```
(* Initial database value and type (schema). *)
val scott1 = db
  : {emps: {name: string, empno: int, deptno: int,
           hiredate: string} bag,
     depts: {deptno: int, name: string} bag};

(* Shim that makes a v1 database look like v2. *)
fun scott2on1shim scott1 =
  {emps =
    fn () => from e in scott1.emps
      yield {e with hiredate = Date.fromString(e.hiredate)},
    depts = fn () => scott1.depts};

(* Shim that makes v3 database look like v1. *)
fun scott1on3shim scott3 =
  {emps =
    fn () => from e in scott3.emps
      yield {e with hiredate = Date.toString(e.hiredate)
            removing rating},
    depts = fn () => scott3.depts};

(* An application writes its queries & views against version 2;
   shims make it work on any actual version. *)
val scott = scott2;
fun recentHires () =
  from e in scott.emps
  where e.hiredate > Date.subtract(Date.now(), 100);
```

# Conclusions





@julianhyde @morel\_lang  
<https://github.com/julianhyde>  
<https://github.com/hydromatic/morel>



Yes, we can do better than SQL!  
Best of both query languages + programming languages  
Give Morel a try!



Morel